

Center for Exascale Radiation Transport

PDT Overview and Development with STAPL

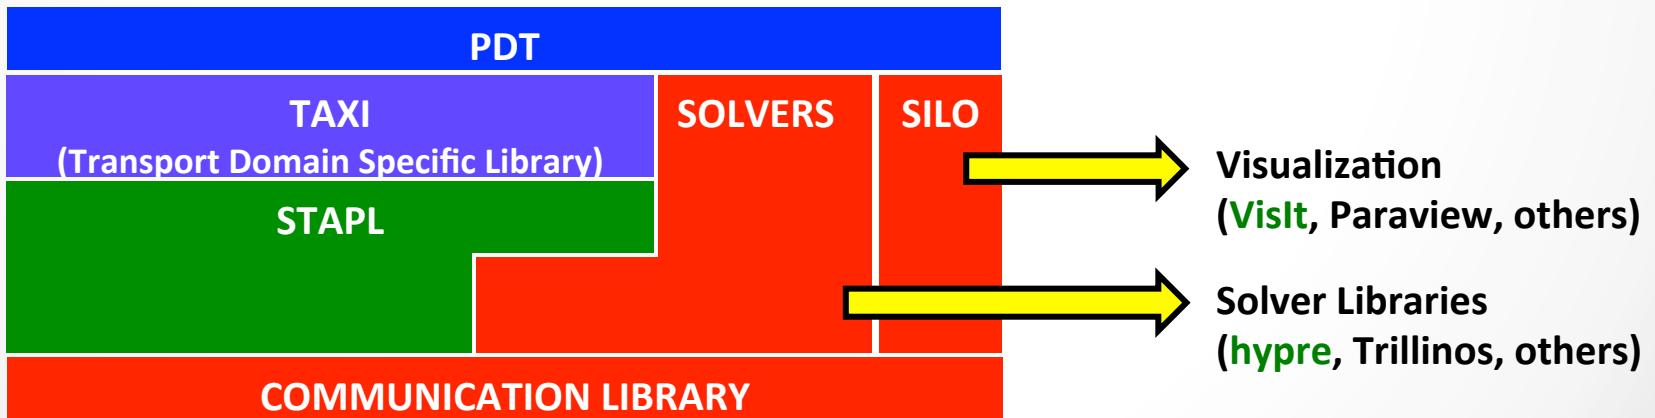
Marvin Adams, Timmie Smith

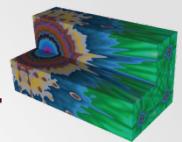
CERT Programming Model Deep Dive

February 18, 2014

PDT

- 2D/3D, massively parallel, deterministic transport code written in C++
- Built on Standard Template Adaptive Parallel Library (STAPL)
- Uses a transport-specific domain library (TAXI)
 - In development (moving code from PDT into TAXI)
 - Specializes STAPL graph/scheduling/partitioning
 - Contains other common code and transport routines
 - Simplifies PDT and allows for rapid prototype of transport algorithms



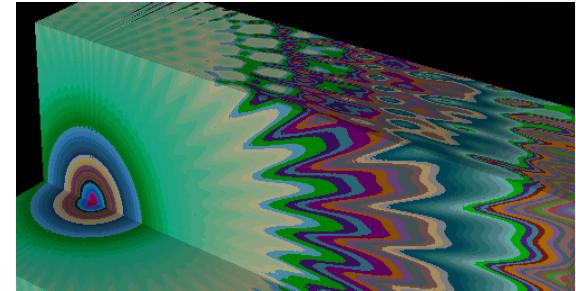
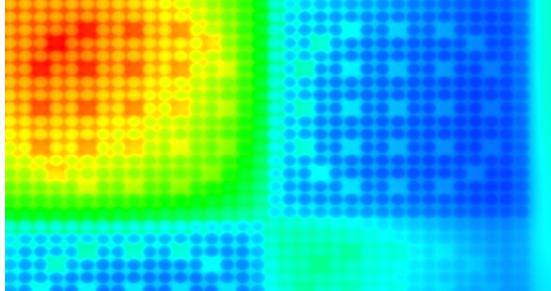
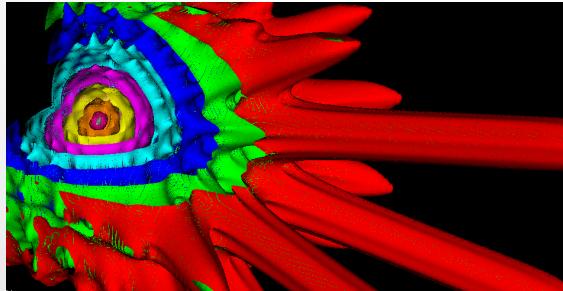


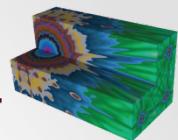
CERT

Center for Exascale Radiation Transport

Capabilities: General

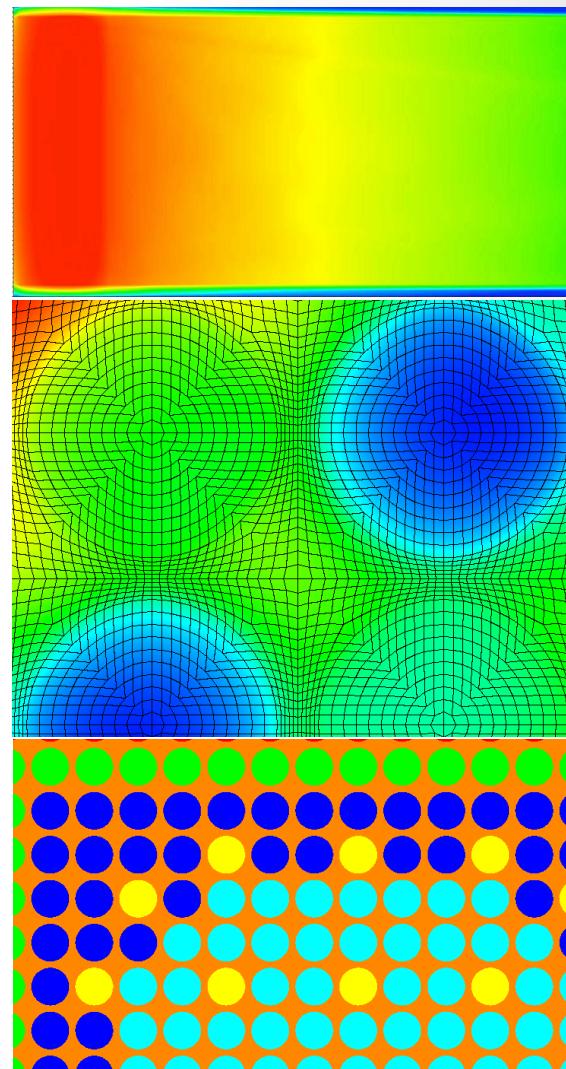
- **Mixed-mode** parallelism (MPI/threads)
(testing and debugging now)
- **Particle types:** neutron, gamma, coupled neutron-gamma, electron, coupled electron-photon, thermal radiation
- **Problem types:** Steady-state, Time-dependent, Criticality, Depletion
- **Uncollided flux:** Novel sophisticated treatment
- Can be called from external programs (DAKOTA, others)

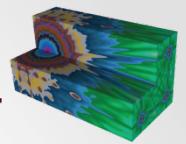




Capabilities: Numerical

- **Discretization:**
 - Energy: Multigroup and FEDS
 - Direction: Discrete ordinates (many quad. types)
 - Spatial: FEM, finite-volume (many options)
 - Time: Implicit (many options, including TBDF2)
 - Anisotropic scattering: arbitrary order
- **Iteration:** sweep-based; many options for accelerating convergence:
 - Diffusion preconditioners (CDFEM and MIP)
 - Gray diffusion for radiative transfer
 - *Can include temperature*
 - Krylov methods
 - Non-linear boundary projection acceleration (*in development*)
 - One-group diffusion for neutron upscattering (*in development*)





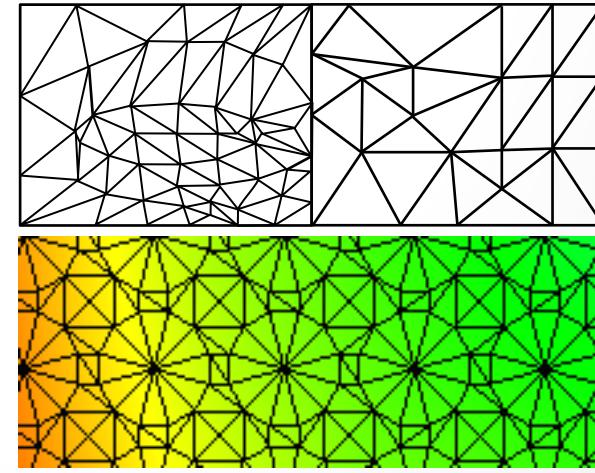
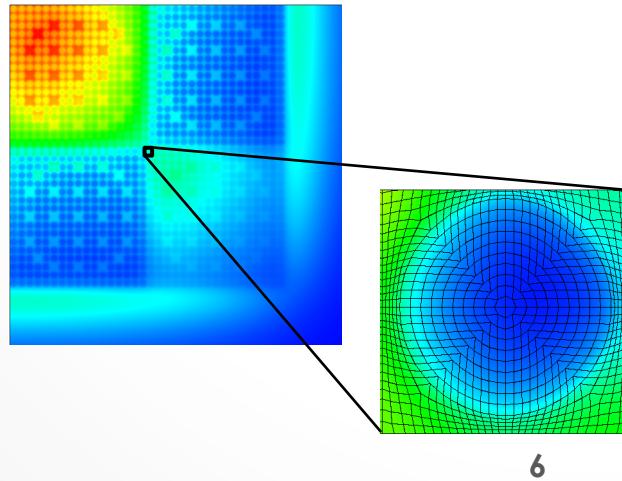
CERT
Center for Exascale Radiation Transport

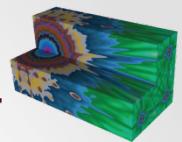
Capabilities: Sweeps

- Scalable parallel transport requires:
 - Scalable parallel algorithm for each iteration
 - Iterations counts that do not increase with resolution
(satisfied by sweep-based methods)
- PDT employs a provably optimal sweep algorithm for rectangular (parallelepiped) *cellsets*:
 - Optimal scheduler that executes sweeps in the fewest possible stages (M&C 2013, M. P. Adams et al.)
 - Performance model that estimates solve time for given problem, machine, partitioning, and aggregation
 - Algorithm to select partitioning and aggregation parameters that minimize estimated solve time
- TAMU/LLNL collaboration has shown that sweeps can scale well to $O(10^6)$ cores

Capabilities: Grids

- PDT has *general* sweep and work functions for arbitrary grids
- PDT has *optimal* sweeps for hierarchical grids: **rectangular** “cellsets” containing **arbitrary** cell shapes & connectivities
 - At cellset level, grid looks brick-like, which permits optimal sweeps
 - Within cellset are many options:
 - Regular rectangular
 - Reactor grids
 - 2D: Triangular (2D), triangular prism (3D), or Tetrahedral (3D) grids built independently and **stitched at interfaces** (forming polygons / polyhedra)



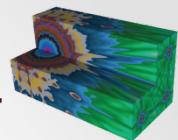


CERT

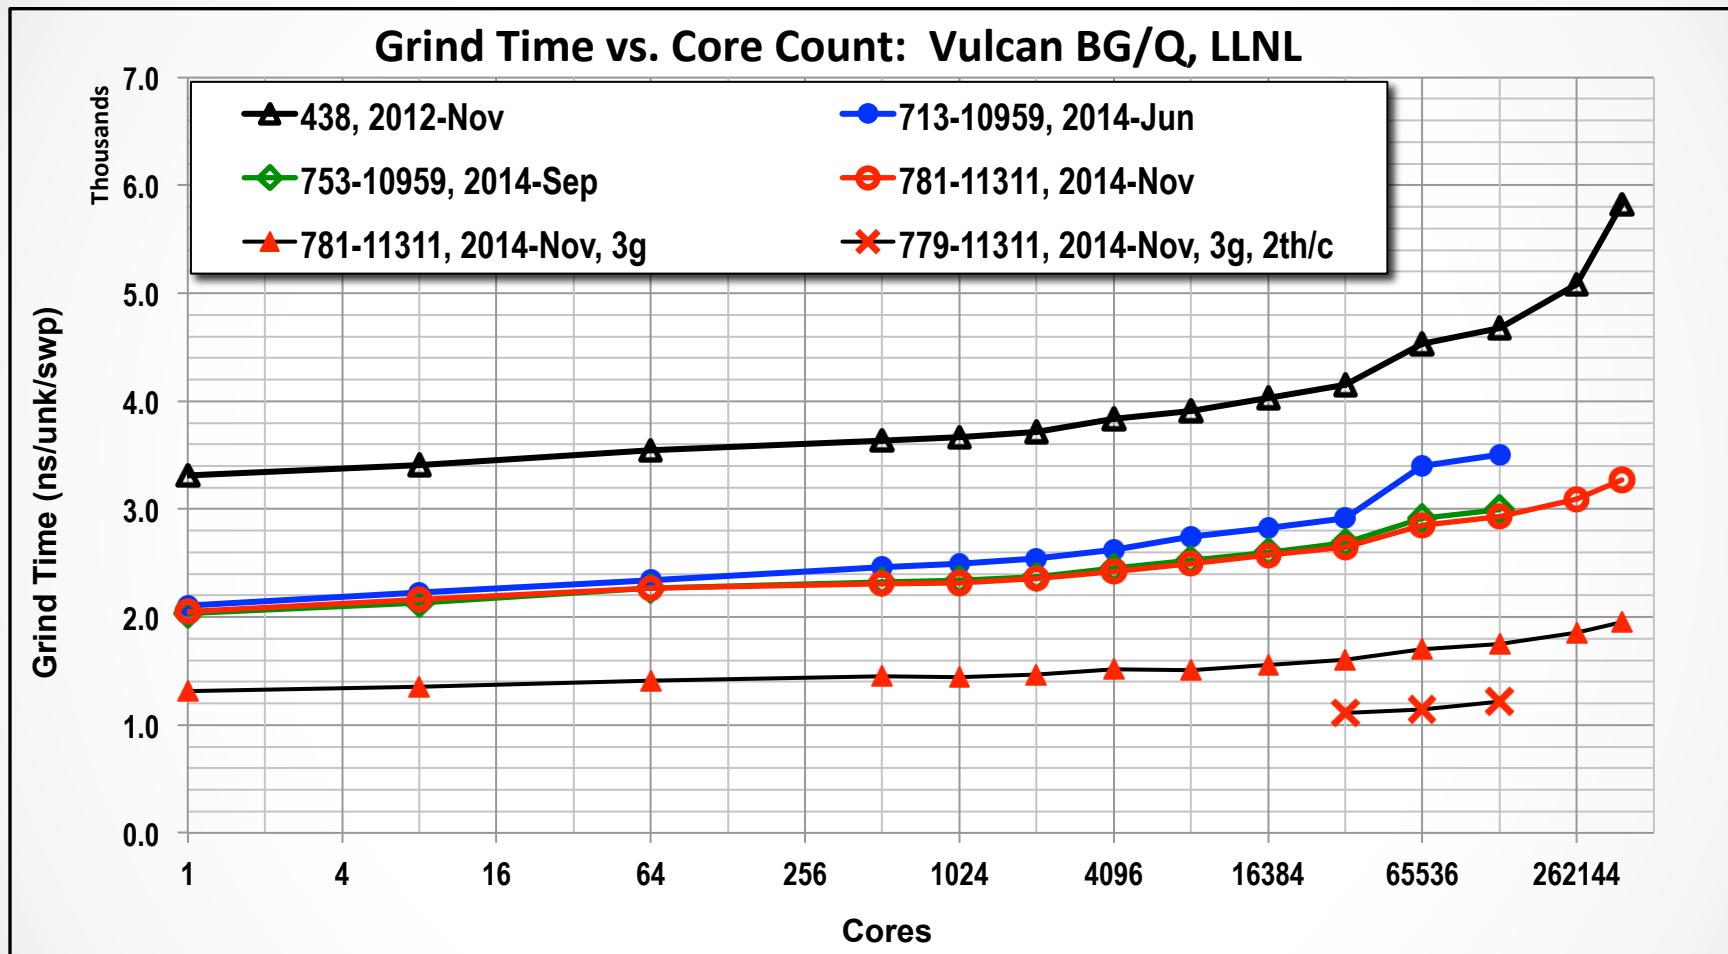
Center for Exascale Radiation Transport

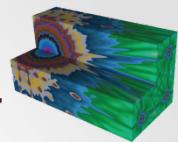
PDT Scaling

- Performance modeling is critical
 - Most helpful in finding code issues and reducing overhead
 - Current model will be extended to mixed-mode and nested parallelism
- Current version of PDT runs with 1M MPI processes
 - Not long ago our largest run was 384K MPI processes
 - Recently completed runs with 512K and 1M MPI processes
 - This used vulcan (BG/Q) at LLNL with 16K nodes and 64 MPI processes/node (one per thread)



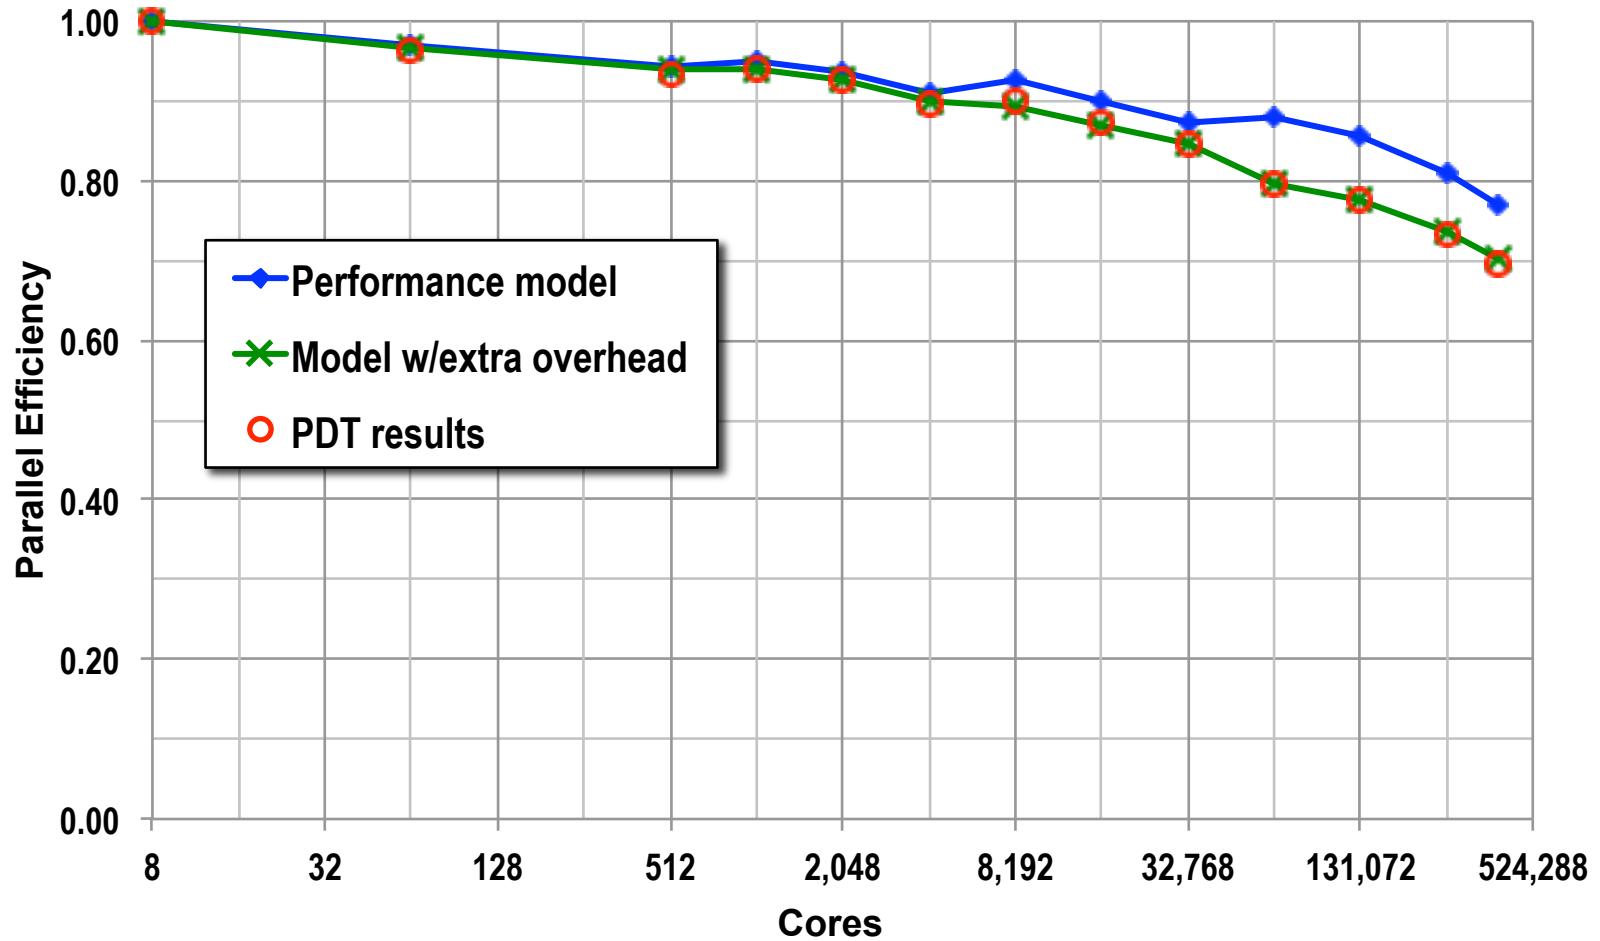
PDT/STAPL Grind Times

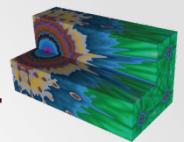




PDT/STAPL Scaling

PDT Parallel Efficiency vs. Core Count, Weak Scaling, IBM BG/Q





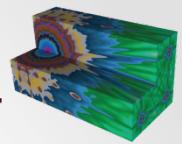
CERT

Center for Exascale Radiation Transport

PDT-STAPL

- Parallelism in PDT is obtained via the underlying STAPL library
- Many of PDT's fundamental data structures are STAPL components

PDT Component	STAPL Component
Grid	graph
Krylov vectors	array
Grid vertices	array
Cellsets	hierarchical_graph_view
Grid cycle detection	strongly_connected_components
Iterative algorithms (inner products, etc.)	inner_product, map, map_reduce
PMPIO functionality	serial_sets skeleton
Task dependencies	explicit paragraphs
Sweep schedule	specialized scheduler

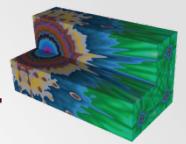


CERT

Center for Exascale Radiation Transport

Integrating STAPL in PDT

- Ease of integrating different components varied
- array and array_view were directly usable in Krylov methods
- Grid and array of vertices (corners of discretized spatial unit)
 - Introduction of use straightforward
 - Data distribution support made construction difficult at the time
- Algorithm and Skeleton use has been simple
 - SCC for cycle detection
 - inner product, map, map_reduce for Krylov methods
 - serial_sets for parallel I/O
 - Skeleton composition to allow merging task graphs is still coming
- PARAGRAPH and runtime use
 - PARAGRAPH construction using explicit task factory a motivating example for Skeletons
 - Strict priority scheduling of sweeps influenced scheduler specialization interface

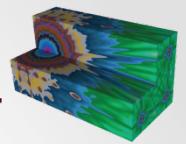


CERT

Center for Exascale Radiation Transport

PDT Grids

- Grid classes for 2D and 3D
- Inherit from `stapl::dynamic_graph`
 - vertex in graph represents unit of discretized spatial domain (cell)
 - edges in graph represent shared faces between cells
- Member functions
 - I/O
 - Vertex motion – skew shape of the cells
 - Boundary initialization
- Data members that are not stored per spatial unit
 - vertices that represent coordinates of corners of cells
 - maximum cell width, other geometry information
- Constructors use cell creator classes
 - Correctly initialize cell
 - Distribution of cells explicitly controlled by creator
 - STAPL data distribution specifications will simplify this

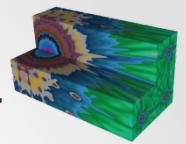


CERT

Center for Exascale Radiation Transport

PDT Sweep Scheduler

- Customized scheduler implemented for executor that processes PARAGRAPHS on each location
 - Choose from set of PARAGRAPHS that contain tasks ready to execute
- Goal: ensure minimum number of scheduling steps
 - Optimal sweep work guarantees that tasks of different sweep directions didn't delay one another
- PARAGRAPHS assigned a priority on each location
 - Intuitively, a PARAGRAPH with more tasks down stream has higher priority
- On each location
 - Process tasks of highest priority PARAGRAPH
 - If highest priority PARAGRAPH doesn't have ready tasks and isn't complete, wait on it
 - After current highest priority PARAGRAPH finishes all tasks on location, proceed to process next priority
- Strict scheduling isn't the default behavior for STAPL
 - Custom scheduler class that implements policy passed to executor at program startup

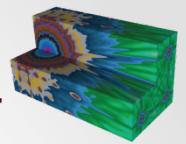


CERT

Center for Exascale Radiation Transport

PDT a Motivating Example

- No STAPL components were custom built for PDT
 - Hierarchical graph_views used in graph partitioning
 - Containers, Algorithms, Skeletons, and PARAGRAPHS used in every STAPL application
- PDT was often the first application to use features
- PDT emphasized need for components
 - Sweep Pattern
 - Hierarchical graph views
 - Skeletons to replace explicit task factories
 - Support for custom data distributions



CERT

Center for Exascale Radiation Transport

PDT and STAPL Co-design

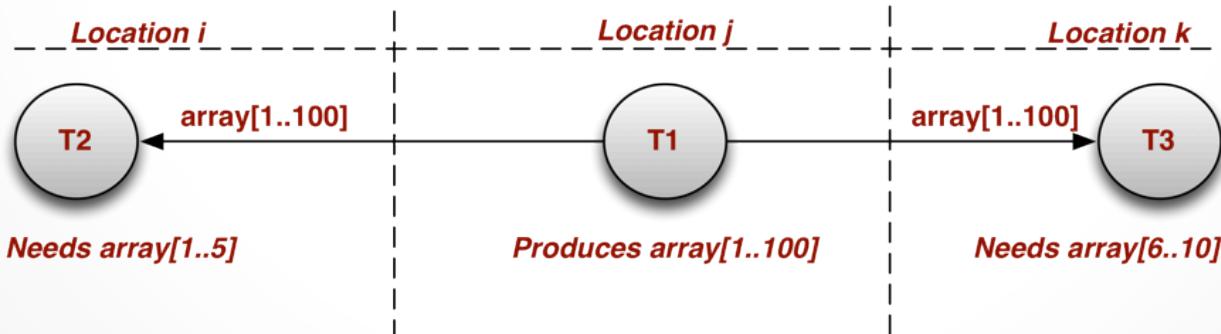
- PDT has driven overhead reduction in STAPL
 - PARAGRAPH and Runtime overheads
 - PDT performance model a lower bound on overhead
- STAPL has simplified development of transport algorithms
 - Fine-grain algorithm specification
 - run-time independent implementation
 - mixed-mode will be introduced without changing code

Partial PARAGRAPH Edge Consumption

Problem

- Coarser tasks increase likelihood successors will only consume to a piece of (aggregate) produced data.
- Finite bandwidth, wasted if we transmit entire edge value.

```
add_task(T1, wf1, v1, 2); // produces array[1..100];
add_task(T2, wf2, consume(T1)); // wf2 filters to [1..5]
add_task(T3, wf3, consume(T1)); // wf3 filters to [6..10]
```



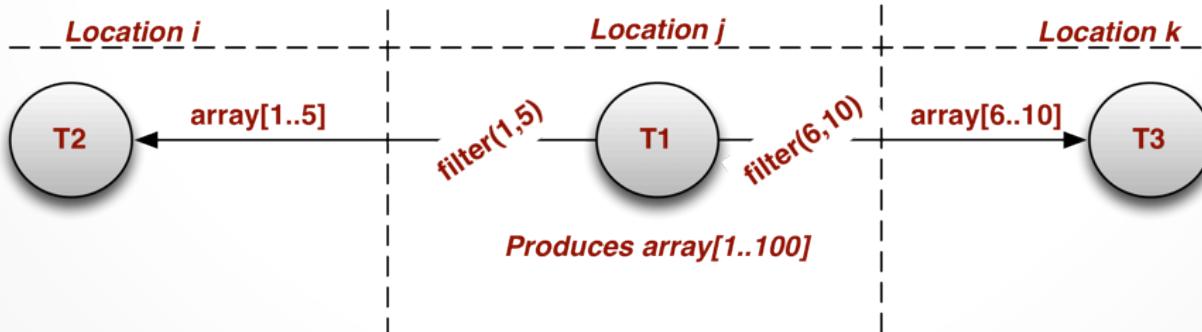
We transport 200 elements, when only 10 are consumed.

Partial PARAGRAPH Edge Consumption

Solution

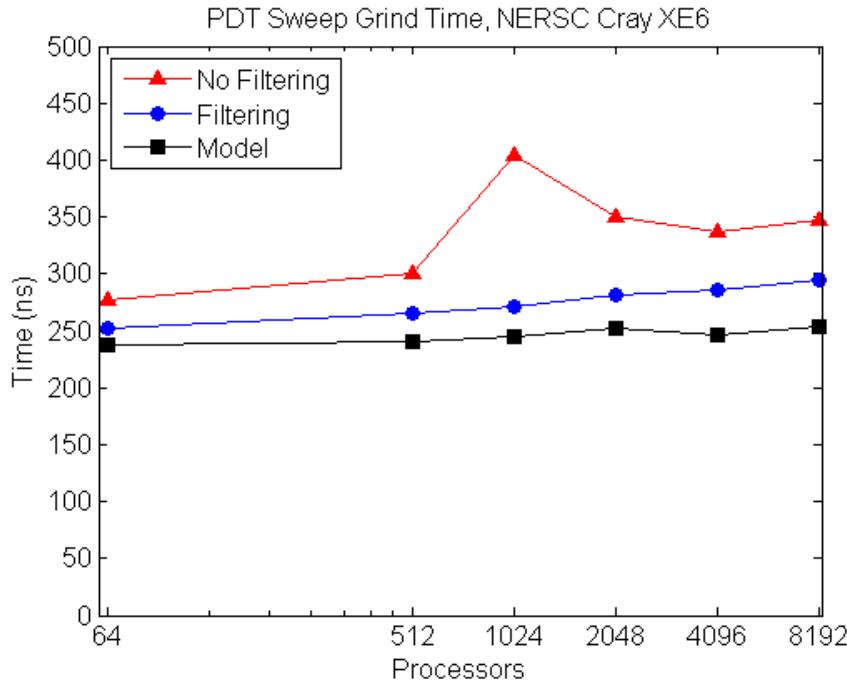
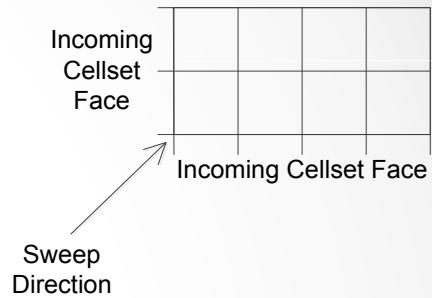
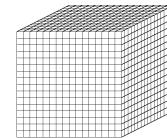
- Filter edge value *prior* to transport.
- Add optional *filtering function* to consume() interface.

```
add_task(T1, wf1, v1, 2); // produces array[1..100];
add_task(T2, wf2, consume(T1, filter_func(1,5));
add_task(T3, wf3, consume(T1, filter_func(6,10));
```



Partial Edge Consumption in PDT

- In sweep, cellsets produce data for **multiple faces** but consume from **one face of each predecessor**.
- Sweep pattern **uses filtering functions**, transparently restrict data prior to calling workfunction.



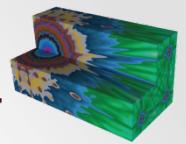
Persistent PARAGRAPHS

Problem

- Portions of global task dependence (e.g., loops) recreate (near) identical **PARAGRAPHS**.
- Even with overlapped creation, redundant & unnecessary computation spent to create **PARAGRAPHS**.

Solution

- Create and execute first **PARAGRAPH** as before, disabling immediate task retirement so that task graph is *persistent*.
- Persistent **PARAGRAPHS** can reused by re-invoking `execute()`.
- Destruction of **PARAGRAPH** object initiates retirement of all tasks and reclamation of resources.



Persistent PARAGRAPHs – Howto

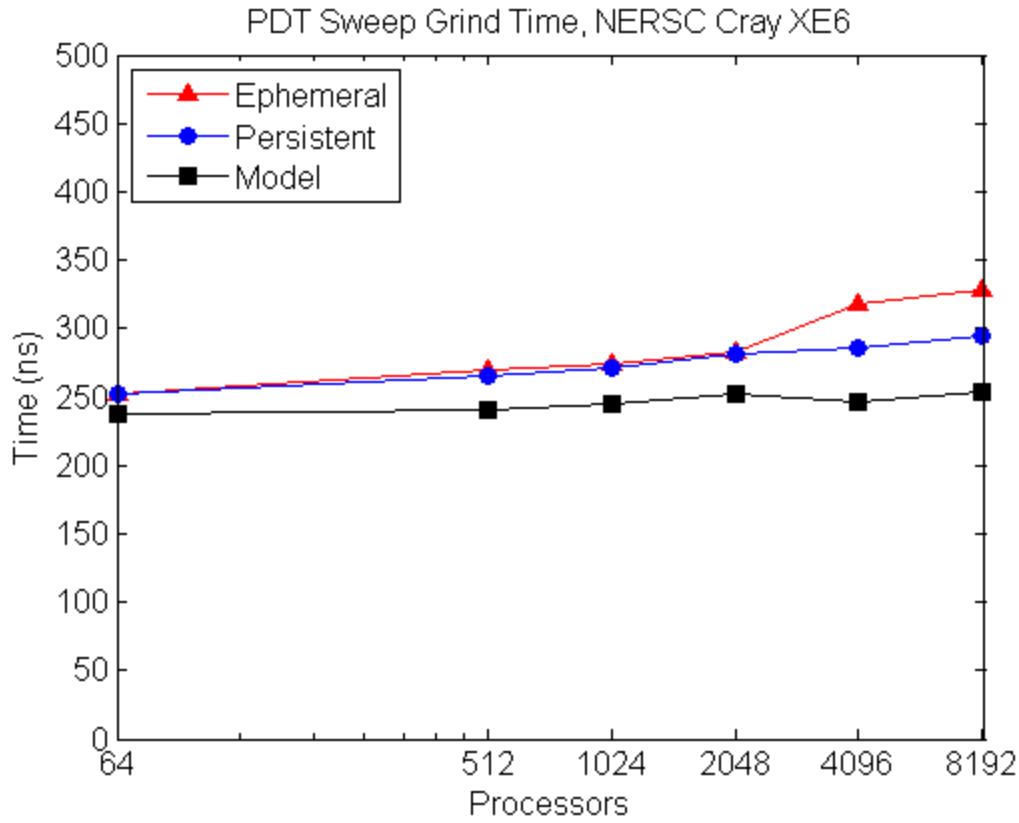
- Enabled by `typedef` flag in the scheduler (per **PARAGRAPH**).
*User guarantees **PARAGRAPH** structure input insensitive.*

```
struct my_scheduler
{
    typedef void enable_persistence;
    ...
};

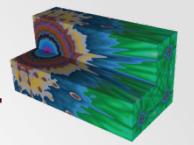
void foo(v1, v2)
{
    pg = make_paragraph(f, v1, v2, my_scheduler());
    retval_1 = pg.execute(); // initialize and execute once.
    ...
    retval_2 = pg.execute(); // reuse for another execution();
}

} // pg destroyed at end of scope. PARAGRAPH dismantled.
```

Persistent PARAGRAPH in PDT Sweep



At 8K, persistency and filtering pushed efficiency from 70% to 78%.

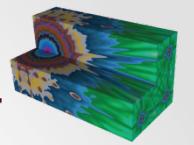


CERT

Center for Exascale Radiation Transport

STAPL simplifies development

- MPI
 - Explicit messaging in user code to express dependencies
 - Manual serialization/deserialization
 - No access to elements on another process
- STAPL
 - Task dependencies enforced by STAPL
 - Serialization only requires method to specify data members
 - Shared object view of containers and views allows simple data access
- Plans to Improve Usability
 - Simplifying or automating generation of define_type method
 - Introduce Skeleton composition in PDT to reduce development of work functions



CERT

Center for Exascale Radiation Transport

Current PDT Development

- Introducing mixed-mode
- Planning introduction of nested parallelism
- Kripke will accelerate process
 - Test support for mixed-mode, mapping on to system
 - Develop approach for nested parallelism in Kripke first
 - Smaller code base with fewer Transport Algorithms to keep online
 - Algorithms developed in Kripke will be applied in PDT

Interoperability in STAPL

Ioannis Papadopoulos
ipapadop@cse.tamu.edu

Why interoperability?



- Existing optimized algorithm implementations
 - STAPL should be able to use them
- Existing applications
 - They should be able to call STAPL effortlessly
- Use views to avoid data copying

STAPL using external libraries



- STAPL can use external libraries
 - BLAS, PBLAS as presented in
Design for Interoperability in STAPL : pMatrices and Linear Algebra Algorithms, A. Buss, T. Smith, G. Tanase, N. Thomas, M. Bianco, N.M. Amato, L. Rauchwerger, *LCPC08*
 - Zoltan redistribution for associative containers (e.g.
`stapl::unordered_map`, `stapl::graph`)
 - STAPL Graph500 benchmark implementation uses Graph500 reference implementation generator
- `stapl::external()` call
 - Invoke non-STAPL distributed memory library functions
 - Halts STAPL; ensures quiescence before function call
- Discussions on how to integrate closer with the `stapl::skeletons` framework

Existing Applications invoking STAPL



- STAPL can be initialized and invoked from existing applications
- PDT has a regular main() program entry point

```
int main(int argc, char* argv[]) {
    MPI_Init_thread(&argc, &argv, MPI_THREAD_SERIALIZED, ...);
    // other MPI code
    auto opt = option{argc, argv} & option{"MPI_Comm", comm};
    stapl::initialize(opts);
    stapl::execute(pdt_entry_point_wf{my_generated_data});
    stapl::finalize();
    // other MPI code
    MPI_Finalize();
    return EXIT_SUCCESS;
}
```

- You can enter and exit STAPL as many times as you want



stapl/benchmarks/kripke

Adam Fidel, Mani Zandifar,
and Nathan Thomas

Parasol Laboratory
Department of Computer Science and Engineering
Texas A&M University





Kripke Overview

- Kripke Overview
- Designing the STAPL implementation
 - Data structures
 - Skeleton based algorithm
 - Views
- Current work



Outline

- An Sn Transport Mini App
- Not a useful calculation - designed to be a representative computational load
- Features all 6 striding orders of directions, groups, and zones
- Specific computational kernel for each “nesting” of domains
- Simplified, no scattering

Kripke Overview

$$\frac{1}{\nu} \frac{\partial \psi}{\partial t} + \Omega \cdot \nabla \psi + \sigma \phi = \sigma_s \psi + \sigma_f \phi + q$$

$$H\Psi^{i+1} = L^+ S L \Psi^i + q$$

$$H\Psi^{i+1} = L^+ I L \Psi^i + 0$$

$$H\Psi^{i+1} = L^+ L \Psi^i$$

- Sweep Kernel (On Core)
 - 3D Diamond Difference
- Parallel Sweep Algo. (MPI)

- LPlusTimes Kernel
(moments -> discrete)

- LTimes Kernel
(discrete -> moments)



Kripke Overview

- Explicit MPI + OpenMP
- Process-level sweep across spatial domain in MPI
- Various kernel configurations
 - With OpenMP, process spatial, energy, and direction domains
 - Primary variables psi / rhs (5D) and sigt (4D).
 - Data layout changes with kernel
 - Layout matches loop nesting
 - Kernel name indicates order of loop nesting (i.e., ZDG -> Zones, Directions, Groups)

Kripke Computation Structure

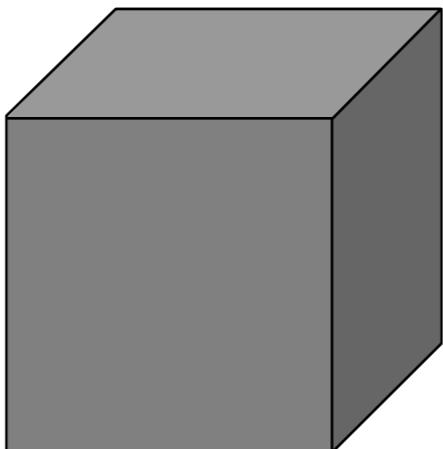
```
Sweep_Solver()
    kernel->LTimes()           - compute phi data

    kernel->LPlusTimes()        - compute rhs data

for (GS)
    forall(DS) {
        MPI_IReceive()          - receive from the upstream of this dir
        if (data is received)
            kernel->Sweep()     - sweep the local grid and compute psi
        MPI_ISend()              - send to the downstream of this dir
    }
```

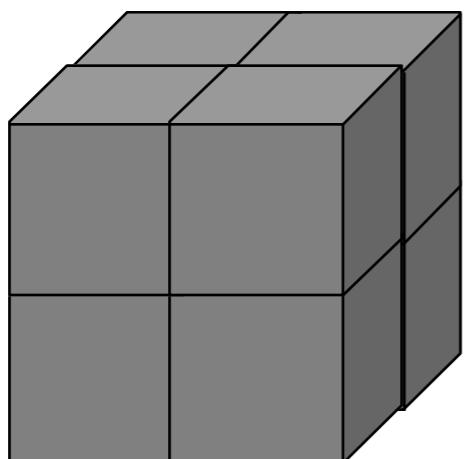
Spatial Decomposition

- Kripke provides spatial decomposition across MPI tasks.
- Each process manages its local grid, which is decided by the number of processors on axis x, y, z.

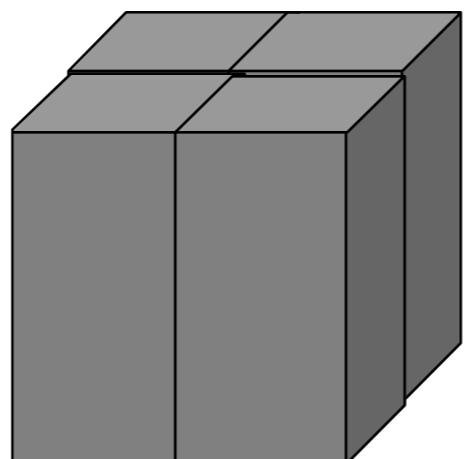


Global space

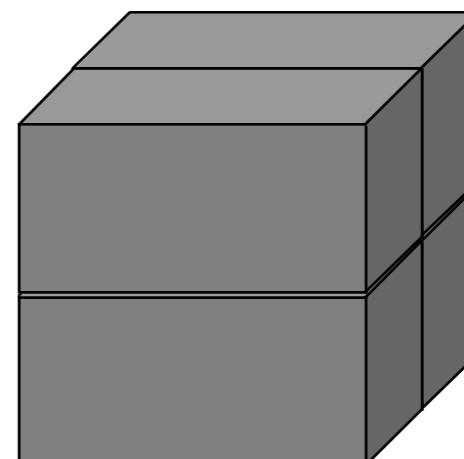
npx : npy : npz
(2 : 2 : 2)



npx : npy : npz
(2 : 1 : 2)

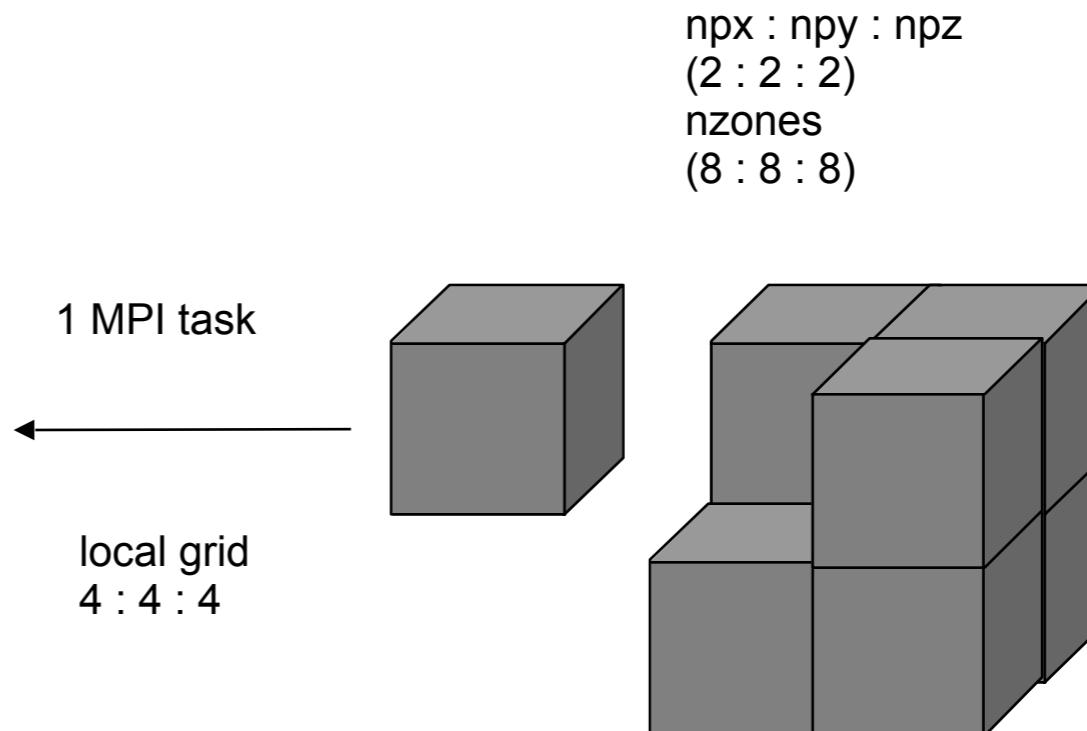
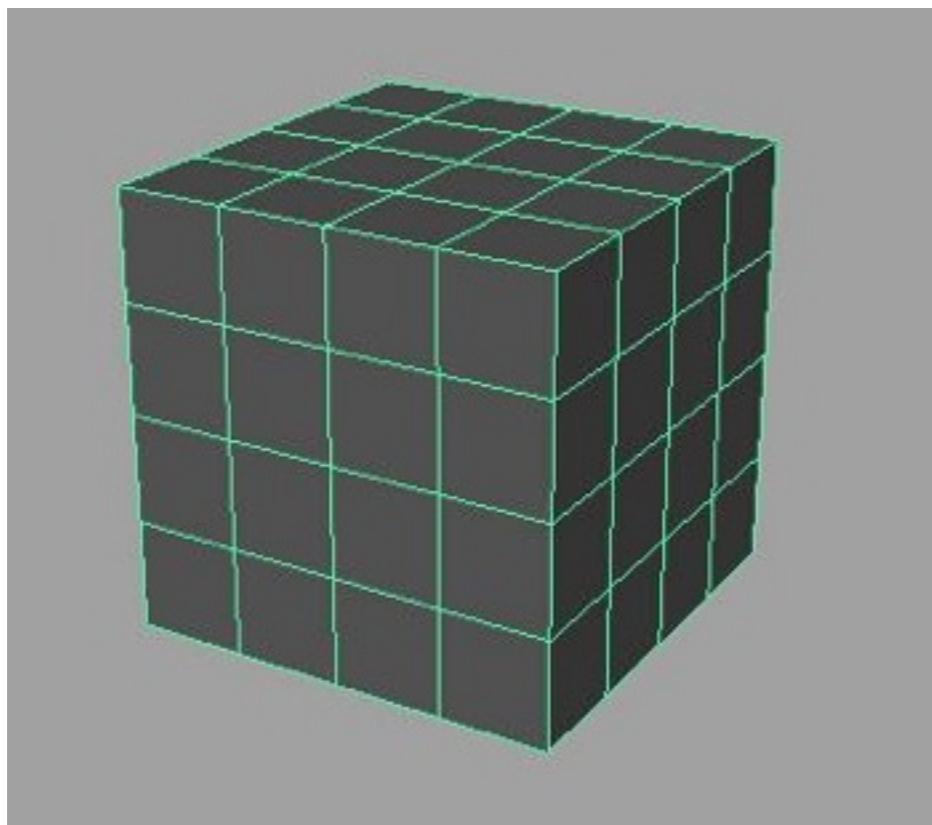


npx : npy : npz
(1 : 2 : 2)



Spatial Decomposition

- One MPI task will covers all the Groups and Directions on a subset of zones.
- Data is transferred through the shared surface of two adjacent MPI tasks.



Parallelism Within Kripke Kernel DGZ

LTimes

```
for each Group Set
  for each Direction Set
    for each entry in L matrix
      for each direction within Direction Set
        forall (groups in set * zones)
```

LPlusTimes

```
for each Group Set
  for each Direction Set
    for each direction within Direction Set
      for each entry in L+ matrix
        forall (groups in set * zones)
```

SweepSolver

```
for each Group Set
```

SweepSolver_GroupSet

```
forall Direction Sets
```

Sweep

```
for each direction within Direction set
  forall groups within Group Set
    for k
      for j | for each zone
        for i |
```

Parallelism Within Kripke Kernel ZGD

LTimes

```
for each Group Set
  for each Direction Set
    forall zones
      for each group within Group Set
        for each direction within Direction Set
          for each entry in L matrix
```

LPlusTimes

```
for each Group Set
  for each Direction Set
    forall zones
      for each group within Group Set
        for each direction within Direction Set
          for each entry in L+ matrix
```

SweepSolver

```
for each Group Set
```

SweepSolver_GroupSet

```
forall Direction Sets
```

Sweep

```
for k | for each zone
  for j |
    for i |
      forall groups within Group set
        for each direction within
          Direction Set
```



Outline

- Kripke Overview
- Designing the STAPL implementation
 - Data structures
 - Skeleton based algorithm
 - Views
- Current work



Implementation Overview

- Unified model for nested parallelism
 - Reuse of algorithms (i.e., same skeleton implementations at both levels of nesting)
 - Addition of support for sweep within sweep
 - ▶ Uses wavefront skeleton
- Abstract machine model, map onto the machine
- Nested containers used to define zonesets.

STAPL Kripke, ZGD Kernel

```
SweepSolver
for each Group Set
    SweepSolver_GroupSet
        forall Direction Sets
            pSweep
                pSweep(subdomain(i,j,k))
                    forall groups within Group set
                        forall each direction within
                            Direction Set
```



Implementation Overview

- Containers
 - Uses multiarray<3, multiarray<5, double>>
 - Mimics setup of Kripke reference code
 - View-based, Composed distribution specifies distribution of both levels
- Algorithms
 - Use compositions of `wavefront` and `zip` skeletons
 - Fine grain solver workfunction
 - [Unit of work:](#)
 - [DD Sweep over One Direction, One Group, and One Zone](#)
- Views
 - Conform data structure to kernel composition
 - New STAPL views - `slices_view` & `extended_view`



Parasol Fine Grain Solver WorkFunction

Unit of work: DD Sweep over One Direction, One Group, and One Zone

```
struct diamond_difference_wf {
    ...
operator()(psi, rhs, sigt, lf_bndry, fr_bndry, bo_bndry)
{
    ...

    double zcos_dzk = 2.0 * m_directions[d].zcos / dzk;
    ... // ycos_dyj, xcos_dxj

    // Calculate new zonal flux
    double new_psi = (rhs + lf_bndry * xcos_dxj
                      + fr_bndry * ycos_dyj
                      + bo_bndry * zcos_dzk) /
                      ((xcos_dxj + ycos_dyj + zcos_dzk + sigt));
    psi = new_psi;

    // Apply diamond-difference relationships
    new_psi *= 2.0;
    lf_bndry = new_psi - lf_bndry;
    ...
}
};
```



Outline

- Kripke Overview
- Designing the STAPL implementation
 - Data structures
 - Skeleton based algorithm
 - Views
- Current work



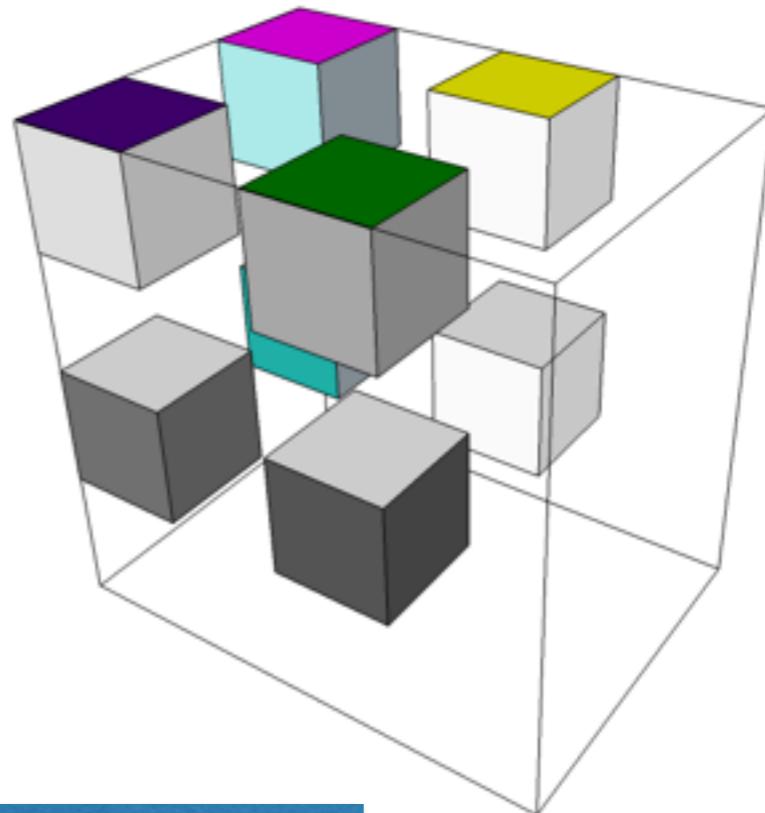
Containers

- Structured spatial decomposition.
 - Use `stapl::multiarray` instead of more general `stapl::graph` (used in PDT).
- Two Level container composition (rhs, psi)
`stapl::multiarray<3, stapl::multiarray<5, double>`
- In contrast to Kripke reference, inner container has an explicit data distribution.
 - Both Distribution and layout changes with kernel
 - Regardless of indexing remains (i, j, k, D, G)

Parasol Zoneset Container

- Outer multiarray has one element per zoneset.
- Uniform distribution in 3D location layout.

```
multiarray<3,...>(  
    stapl::uniform_nd({2, 2, 2}, {4, 4, 4})  
);
```



2x2x2 Zoneset Layout

4x4x4 Location Layout



Kernel Container

- Explicit data distribution as compared to shared memory, OpenMP reference
 - Balanced distribution of Spatial, Direction, or Group Domain (whichever is on the outside).
 - Constant indexing (i, j, k, D, G)
 - Changing data layout (traversal)

DGZ

```
typedef ints_<3>           part_dims;
typedef ints_<3,4,0,1,2>    traversal_t;
multiarrray<5>(
    partial_balanced<part_dims, traversal_t>({8,8,8,4,2})
);
```



Kernel Container Example

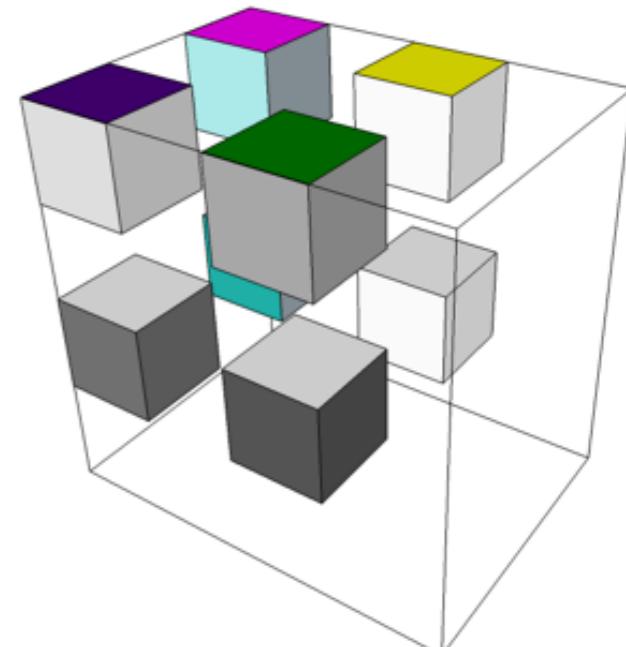
ZDG

```
typedef ints_<0,1,2>          part_dims;
typedef ints_<0,1,2,3,4>    traversal_t;
multiarray<5, ...> ma(
    partial_balanced<part_dims, traversal_t>({8,8,8,4,2})
);
```



Parasol Composing the Distributions

- Previous examples showed inner and outer containers made in isolation
- Actually composed container instantiation made in single statement:
multiarrray<multiarrray<...>> psi (...) ;
- Need to combine distribution specifications into single, composed distribution (view-based)
- Address mapping of each inner container to subset of locations





Composing the Distributions

```
composed_spec vw(...);  
multiarray<multiarray<...>> ma(vw);
```

```
composed_spec_wf {  
    // Specify Outer Container Distribution  
    operator() () {  
        return uniform_nd(...);  
    }  
  
    // Specify Inner Container Distribution  
    operator() (typename Spec1::gid_type gid) {  
        location_range = f(gid);  
        return partial_balanced<...>(size, location_range);  
    }  
};
```

- Kripke Overview
- Designing the STAPL implementation
 - Data structures
 - Skeleton based algorithms
 - Views
- Current work



Kripke Algorithm

Reference Implementation

```
void SweepSolver (User_Data *user_data)
{
    Kernel *kernel = user_data->kernel;
    Grid_Data *grid_data = user_data->grid_data;

    for(int iter = 0;iter < user_data->niter;++ iter){
        // Discrete to Moments transformation
        kernel->LTimes(grid_data);

        // Moments to Discrete transformation
        kernel->LPlusTimes(grid_data);

        for (int group_set = 0;group_set < num_group_sets;++ group_set) {
            doall(int direction_set = 0; ds < num_direction_sets; direction_set++) {
                ...
                kernel->sweep(grid_data,
                                dir_sets[direction_set],
                                i_plane_data[direction_set],
                                j_plane_data[direction_set],
                                k_plane_data[direction_set]);
                ...
            }
        }
    }
}
```



Parasol Sweep Kernels

- Sweep kernel is specialized over the loop orders
 - Zone-Group-Direction (ZGD)
 - Zone-Direction-Group (ZDG)
 - Group-Zone-Direction (GZD)
 - Group-Direction-Zone (GDZ)
 - Direction-Zone-Group (DZG)
 - Direction-Group-Zone (DGZ)
- The loop orders are interchanged



Reference Implementation

ZDG vs. ZGD

```
Kernel_ZGD::sweep(args...)
{
    for (int k = extent.start_k; k != extent.end_k; k += extent.inc_k)      // #1
        for (int j = extent.start_j; j != extent.end_j; j += extent.inc_j)      // .
            for (int i = extent.start_i; i != extent.end_i; i += extent.inc_i) // .
                for (int group = 0; group < num_groups; ++group)           // #2
                    for (int d = 0; d < num_directions; ++d)                  // #3
                        diamond_difference_wf(...);
}
```

```
Kernel_ZDG::sweep(args...)
{
    for (int k = extent.start_k; k != extent.end_k; k += extent.inc_k)      // #1
        for (int j = extent.start_j; j != extent.end_j; j += extent.inc_j)      // .
            for (int i = extent.start_i; i != extent.end_i; i += extent.inc_i) // .
                for (int d = 0; d < num_directions; ++d)                      // #3
                    for (int group = 0; group < num_groups; ++group)          // #2
                        diamond_difference_wf(...);
}
```





Parasol Zip Skeletons

- We use the zip skeleton for loops #2 and #3

```
Kernel_ZGD::sweep(args...)
{
    for (int k = extent.start_k; k != extent.end_k; k += extent.inc_k)      // #1
        for (int j = extent.start_j; j != extent.end_j; j += extent.inc_j)      // .
            for (int i = extent.start_i; i != extent.end_i; i += extent.inc_i) // .
                for (int group = 0; group < num_groups; ++group)           // #2
                    for (int d = 0; d < num_directions; ++d)                  // #3
                        diamond_difference_wf(...);
}
```

```
Kernel_ZGD::sweep(args...)
{
    for (int k = extent.start_k; k != extent.end_k; k += extent.inc_k)      // #1
        for (int j = extent.start_j; j != extent.end_j; j += extent.inc_j)      // .
            for (int i = extent.start_i; i != extent.end_i; i += extent.inc_i) // .
                exec.execute(zip(zip(diamond_difference_wf()))),
                                input_views...
}
```



Parasol Wavefront Skeleton

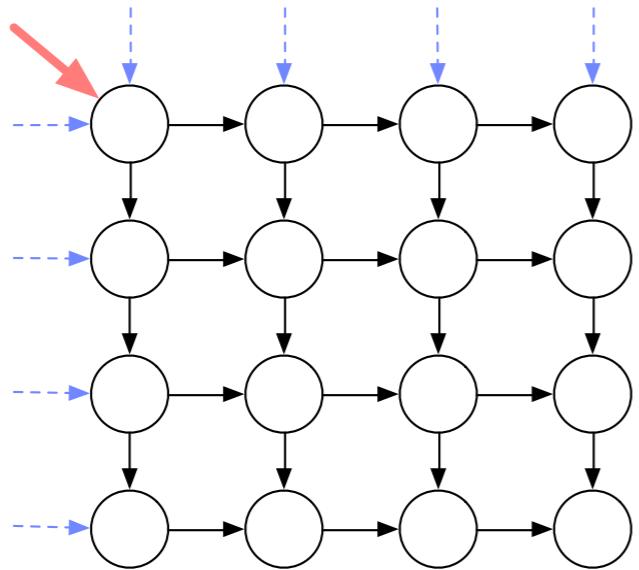
- We use the wavefront skeleton for loop #1

```
Kernel_ZGD::sweep(args...)
{
    for (int k = extent.start_k; k != extent.end_k; k += extent.inc_k)      // #
        for (int j = extent.start_j; j != extent.end_j; j += extent.inc_j)      // .
            for (int i = extent.start_i; i != extent.end_i; i += extent.inc_i) // .
                for (int group = 0; group < num_groups; ++group)           // #
                    for (int d = 0; d < num_directions; ++d)                  // #
                        diamond_difference_wf(...);
}
```

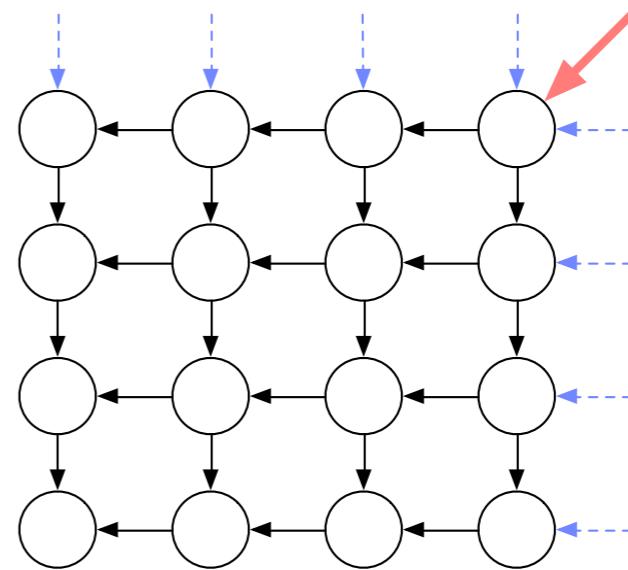
-
- A topologically ordered traversal of a data structure
 - Wavefront on regular grids in STAPL
 - 2D wavefront skeleton
 - 3D wavefront skeleton

Parasol 2D Wavefront Skeleton

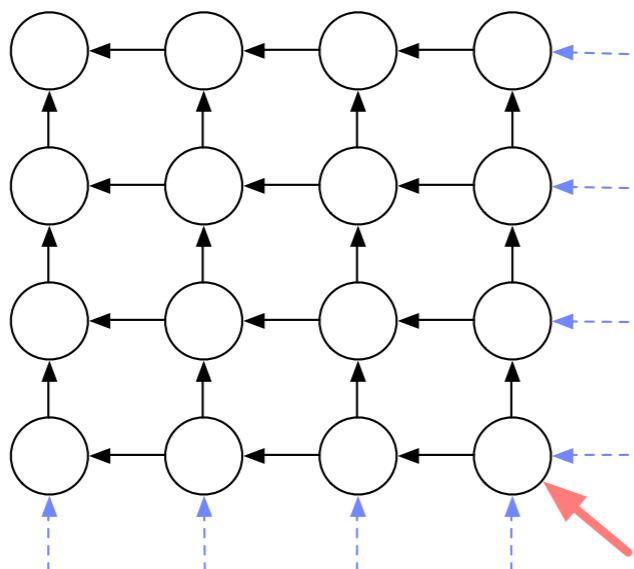
- Parameterized over the size and corner to start from



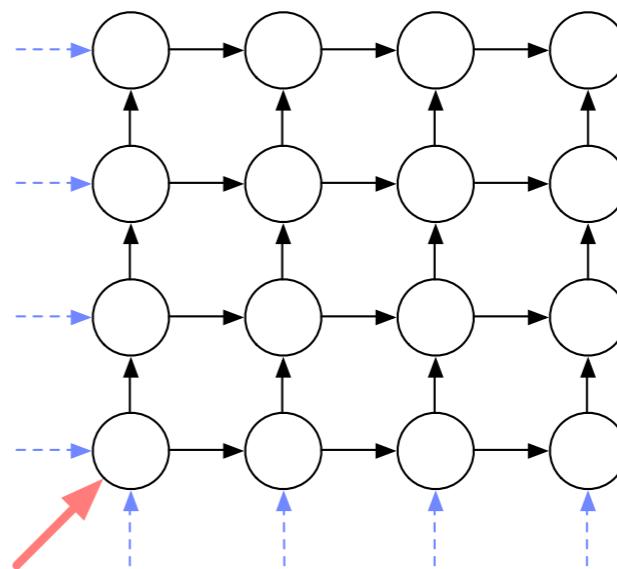
Starting from top left ({first, first})



Starting from top right ({first, last})



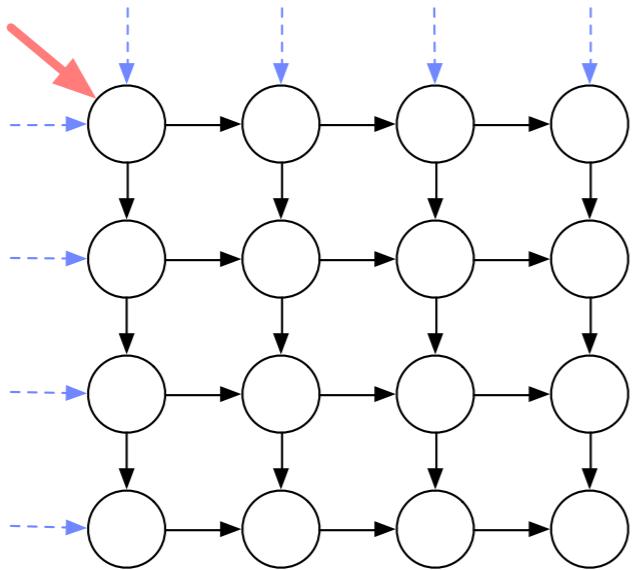
Starting from bottom right ({last, last})



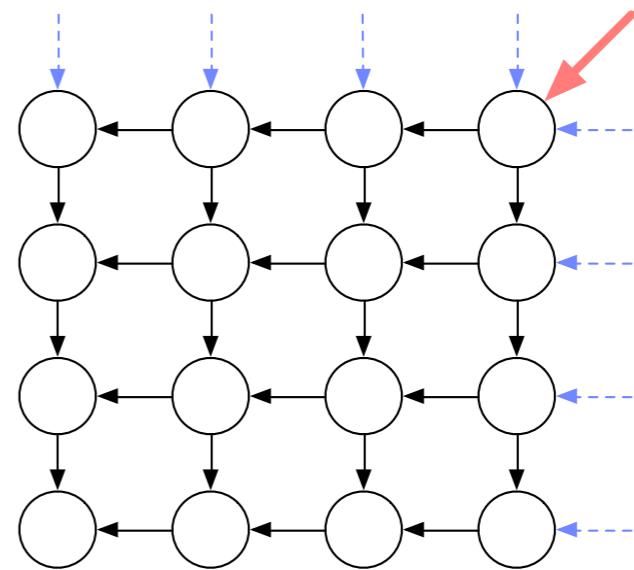
Starting from bottom left ({last, first})

Parasol 2D Wavefront Skeleton

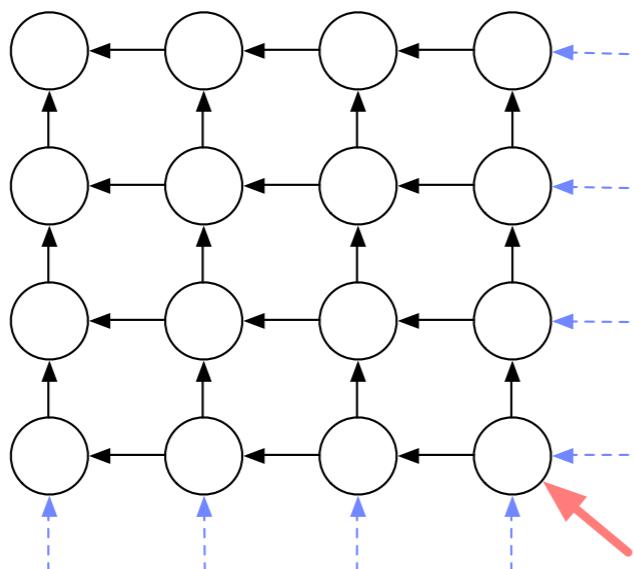
- In STAPL they are defined as:



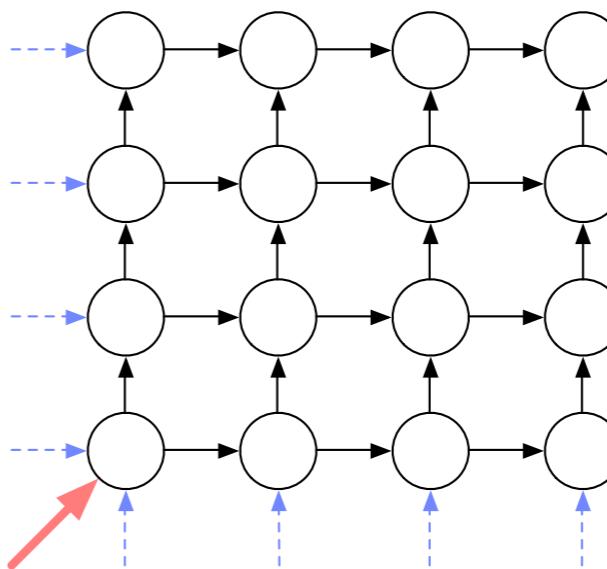
`skeletons::wavefront(op, {{first, first}})`



`skeletons::wavefront(op, {{first, last}})`



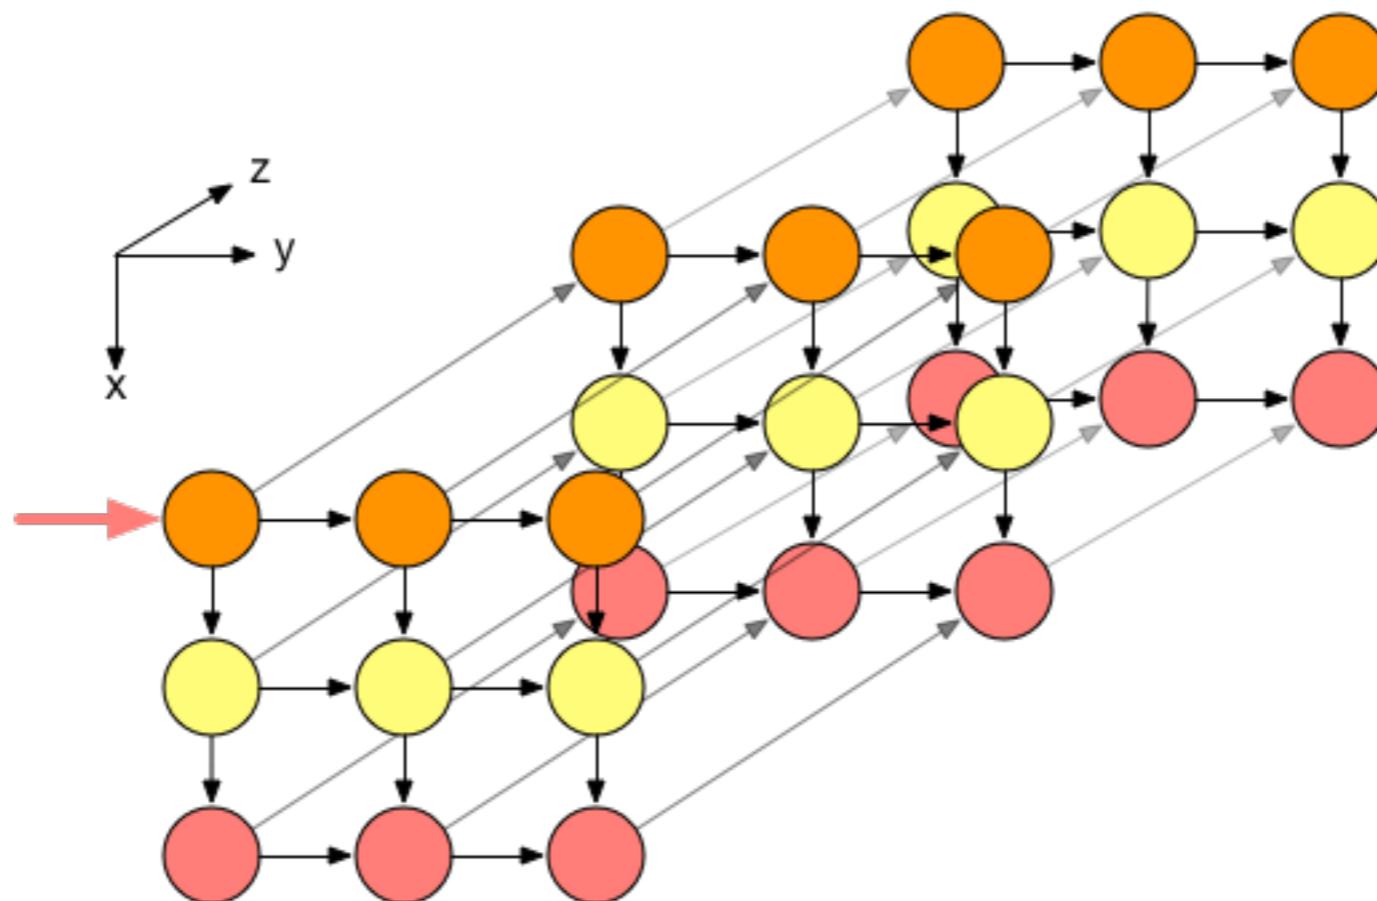
`skeletons::wavefront(op, {{last, last}})`



`skeletons::wavefront(op, {{last, first}})`

Parasol 3D Wavefront Skeleton

- Defined similarly
- Parameterized based on the corner to start from



`wavefront(op, {{first, first, first}})`



Parasol 3D Wavefront Skeleton Example

- Implementation in STAPL

```
// Definition in the STAPL Skeleton Framework
template <typename Op, int i>
auto wavefront(Op op, std::array<wavefront_direction, i> corners) {
    return elem<input_span, doacross>(wavefront_pd(op, corners));
}

// Assuming the input is provided in the input_view, you need to
// create your boundary views
auto size = input_view.size();
auto boundary_view1 = functor_view(size, your_generator(direction0));
auto boundary_view2 = functor_view(size, your_generator(direction1));
auto boundary_view3 = functor_view(size, your_generator(direction2));

// Finally, execute the wavefront skeleton
algorithm_executor exec;
exec.execute(wavefront(my_wavefront_operation(),
                     {{first, first, first}}),
            boundary_view1, boundary_view2, boundary_view3,
            input_view);
```

- Reference Implementation

```
Kernel_ZGD::sweep(args...)
{
    for (int k = extent.start_k; k != extent.end_k; k += extent.inc_k)      // #1
        for (int j = extent.start_j; j != extent.end_j; j += extent.inc_j)      // .
            for (int i = extent.start_i; i != extent.end_i; i += extent.inc_i) // .
                for (int group = 0; group < num_groups; ++group)           // #2
                    for (int d = 0; d < num_directions; ++d)
                        diamond_difference_wf(...);
}
```

- STAPL Implementation

```
Kernel_ZGD::sweep(args...)
{
    auto kernel_ZGD_skeleton = waveform(zip(zip(diamond_difference_wf(...)),
                                              find_corners<ZGD>(...)));
    exec.execute(kernel_skeleton, input_views...);
}
```



Parasol All Kernels

```
auto kernel_ZGD_skeleton = waveform(zip(zip(diamond_difference_wf(...)),  
                                         find_corners<ZGD>(...)));  
  
auto kernel_ZDG_skeleton = waveform(zip(zip(diamond_difference_wf(...)),  
                                         find_corners<ZGD>(...)));  
  
auto kernel_DZG_skeleton = zip(waveform(zip(diamond_difference_wf(...)),  
                                         find_corners<ZGD>(...)));  
  
auto kernel_DGZ_skeleton = zip(zip(waveform(diamond_difference_wf(...),  
                                         find_corners<ZGD>(...))));  
  
auto kernel_GZD_skeleton = zip(waveform(zip(diamond_difference_wf(...)),  
                                         find_corners<ZGD>(...)));  
  
auto kernel_GDZ_skeleton = zip(zip(waveform(diamond_difference_wf(...),  
                                         find_corners<ZGD>(...))));
```



Parasol Kripke Algorithm

- The full algorithm is defined as

```
// kernel computation
auto kernel_ZGD_skeleton = waveform(zip(zip(diamond_difference_wf(...)),
                                         find_corners<ZGD>(...)));
for (int iteration = 0; iteration < num_iterations; ++iteration) {
    exec.execute(compose(zip(LTimes(...)), zip(LPlusTimes(...))),
                 input_view...);

    for (int group_set = 0; group_set < num_group_sets; ++group_set) {
        exec.execute(doall(kernel_ZGD_skeleton,
                           [&](size_t size){return num_direction_sets}),
                     input_view...)
    }
}
```



- Kripke Overview
- Designing the STAPL implementation
 - Data structures
 - Skeleton based algorithms
 - [Views](#)
- Current work



Algorithms and Views

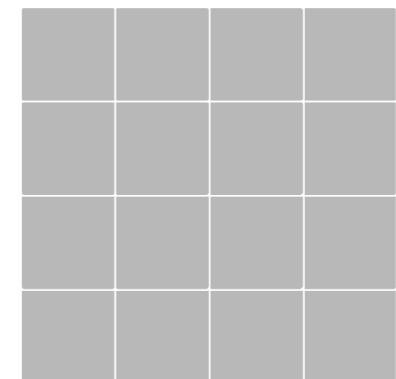
- Data structure inside zoneset is 5-dimensional
 - Separate 3 dimensions for waveform
 - Separate others into 1D zips
- Modify how a container looks to conform with what the algorithm expects
 - View the 5D container in a manner conformant to the kernel nest
 - Auxiliary 4D structure (sigt) where direction dimension is aggregated



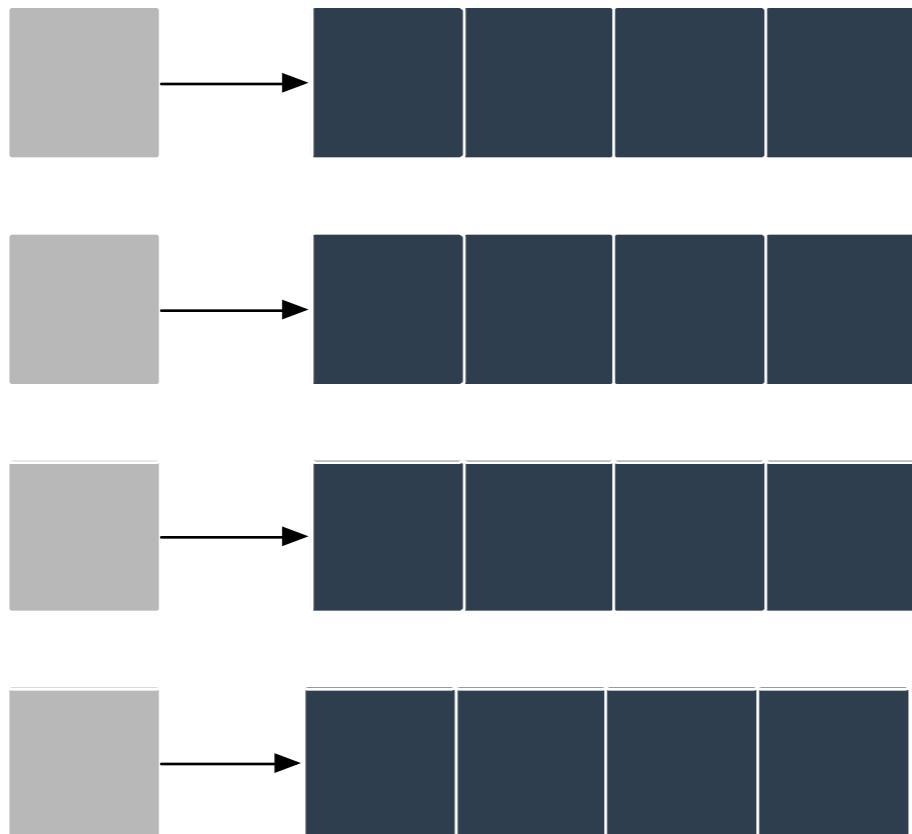
Multiaray View

- The main container in Kripke is `stapl::multiaray`
- `multiaray_view` is a view over the multiaray container
- Forms the basis of the following complex views
 - Slices view
 - Extended view
 - Functor view

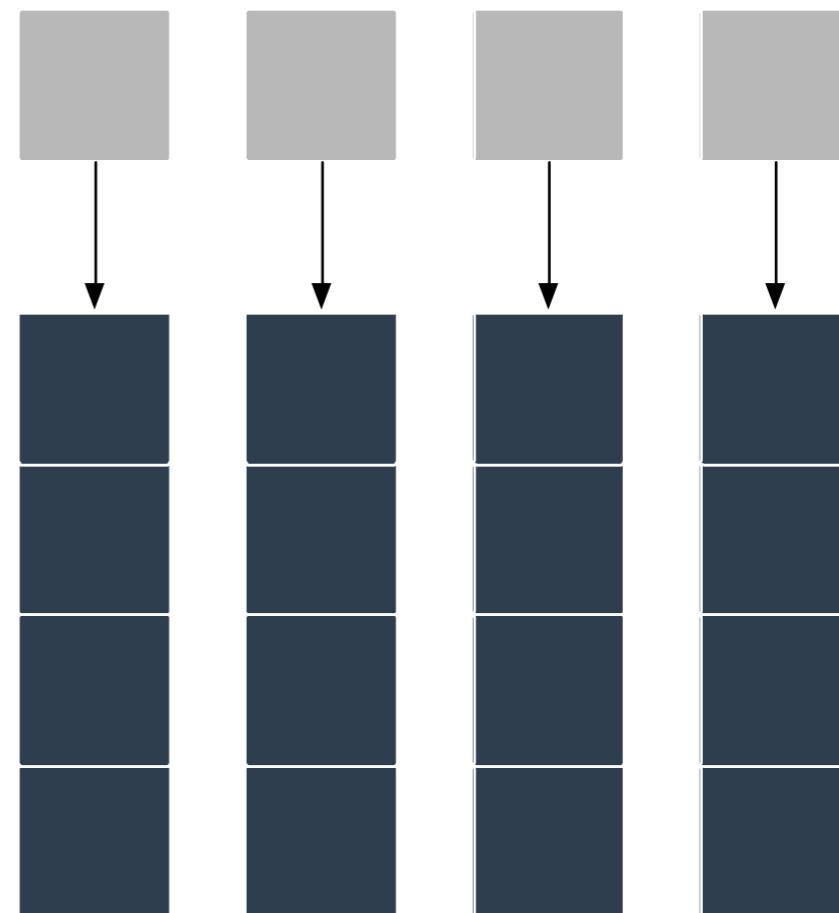
- For a 2D matrix, produce a collection whose elements are rows or columns



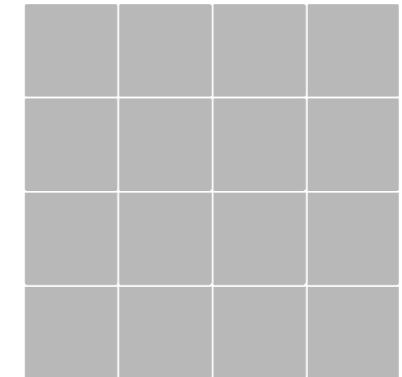
rows(A)



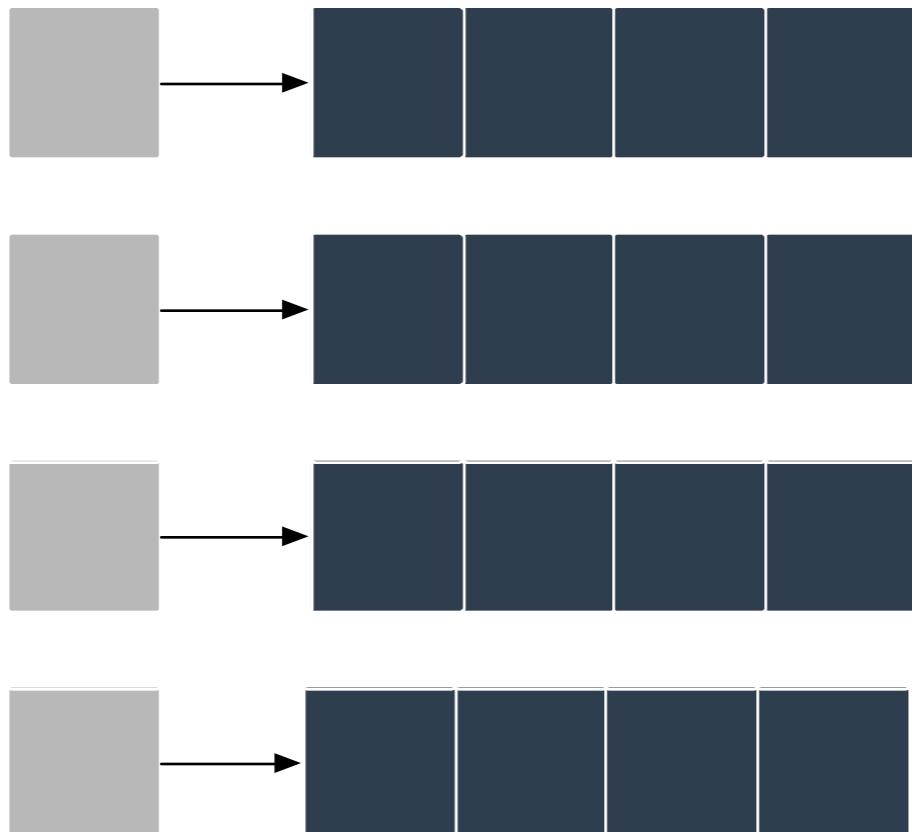
columns(A)



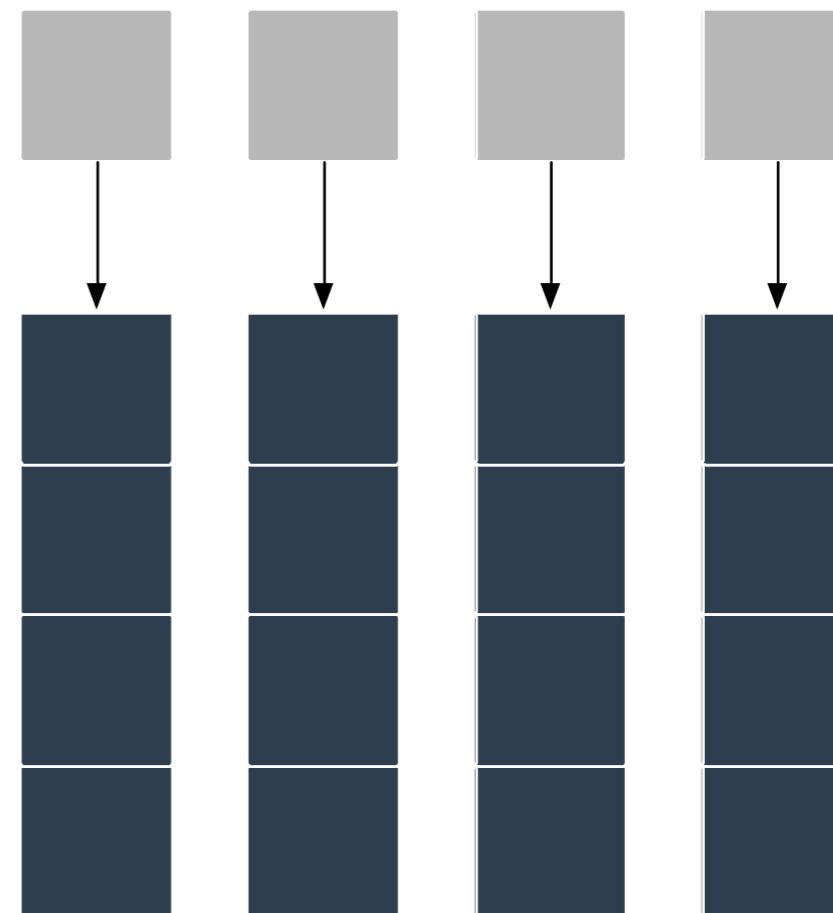
- View that produce subviews
 - Each subview is a single slice along given dimension

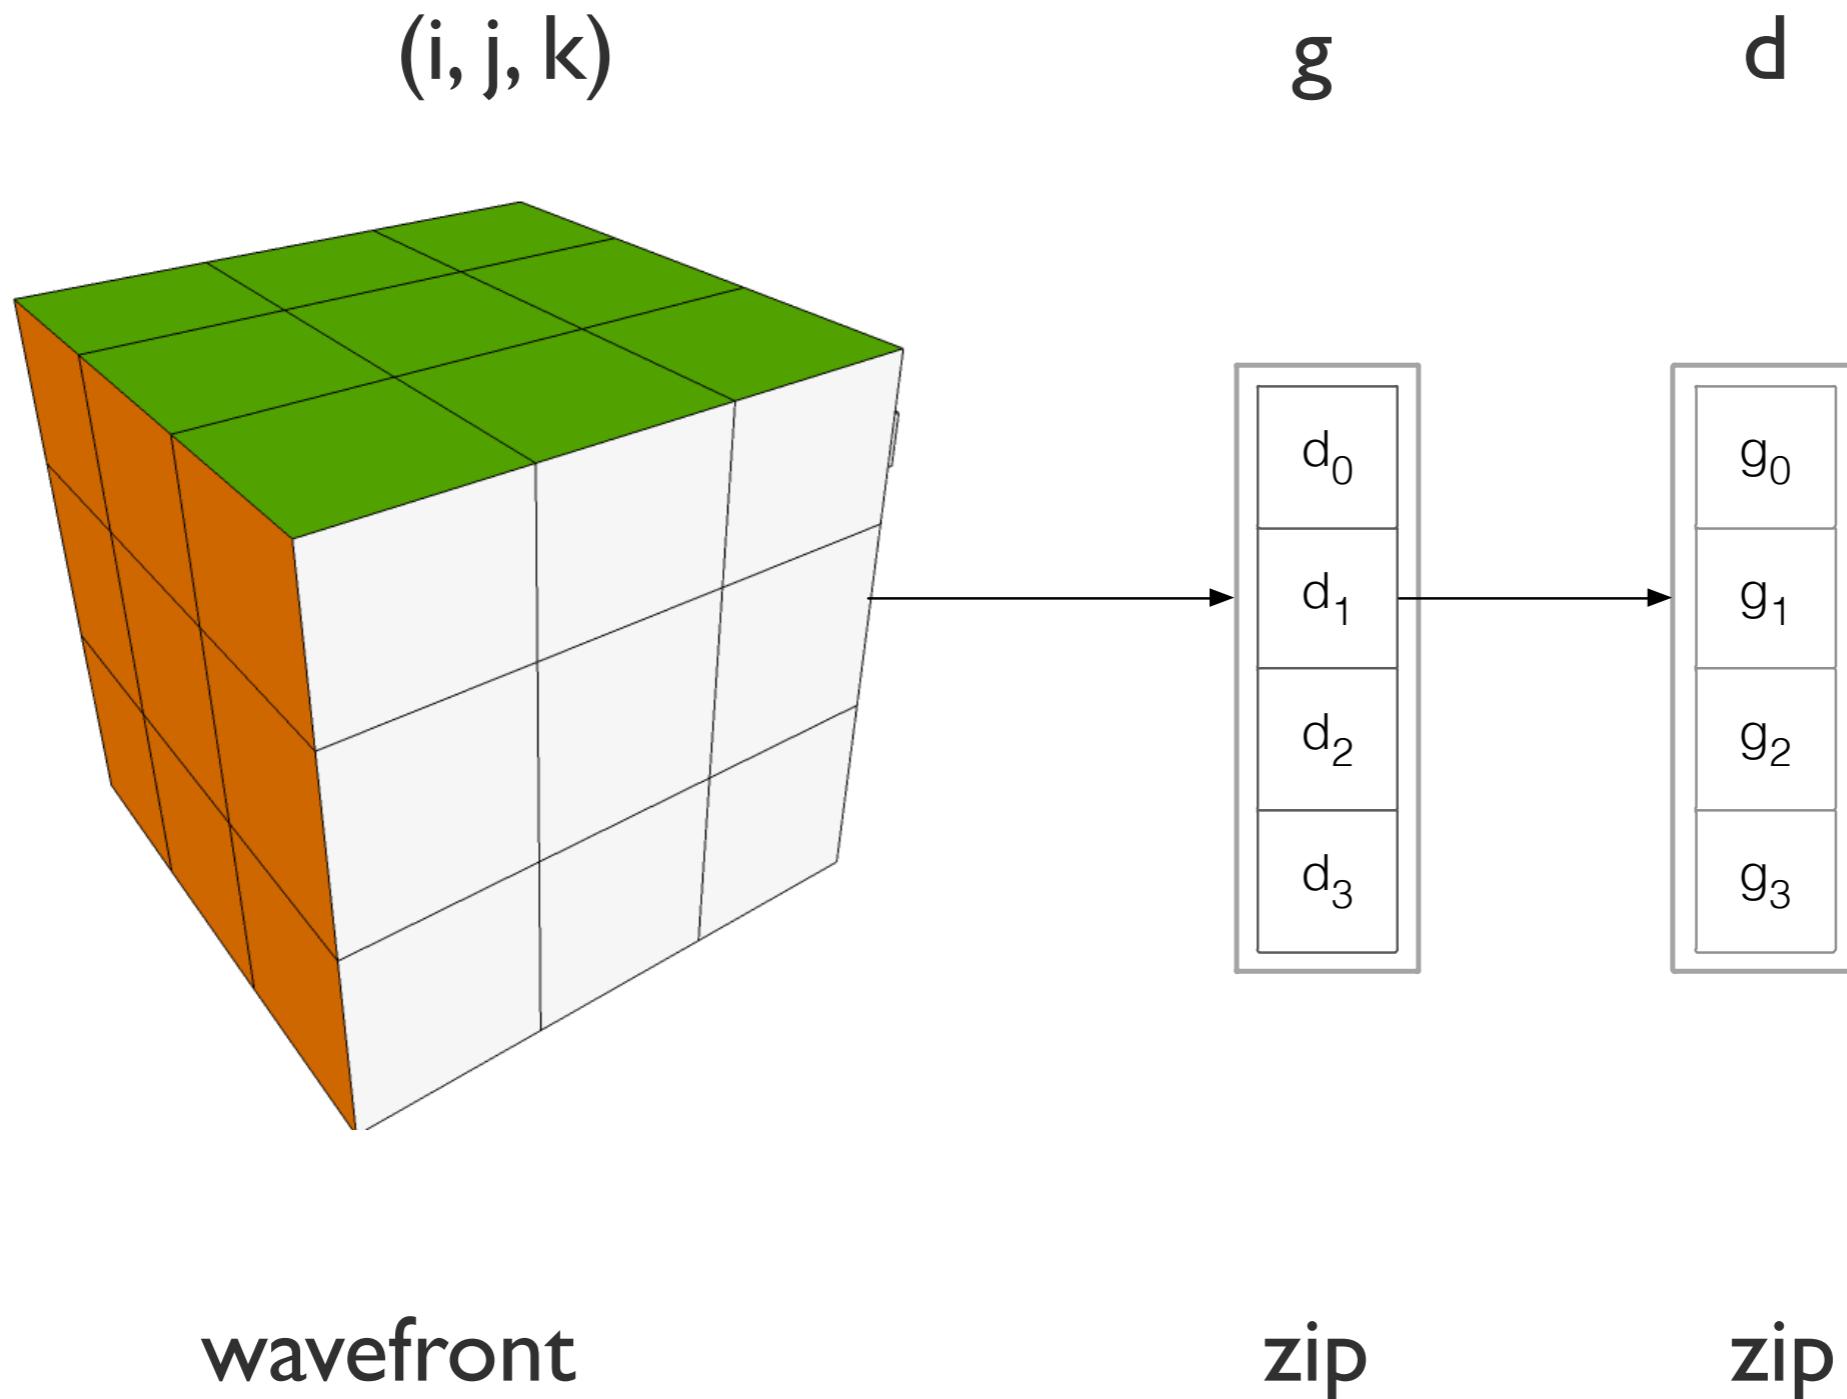


`make_slices_view<0>(vw)`



`make_slices_view<1>(vw)`



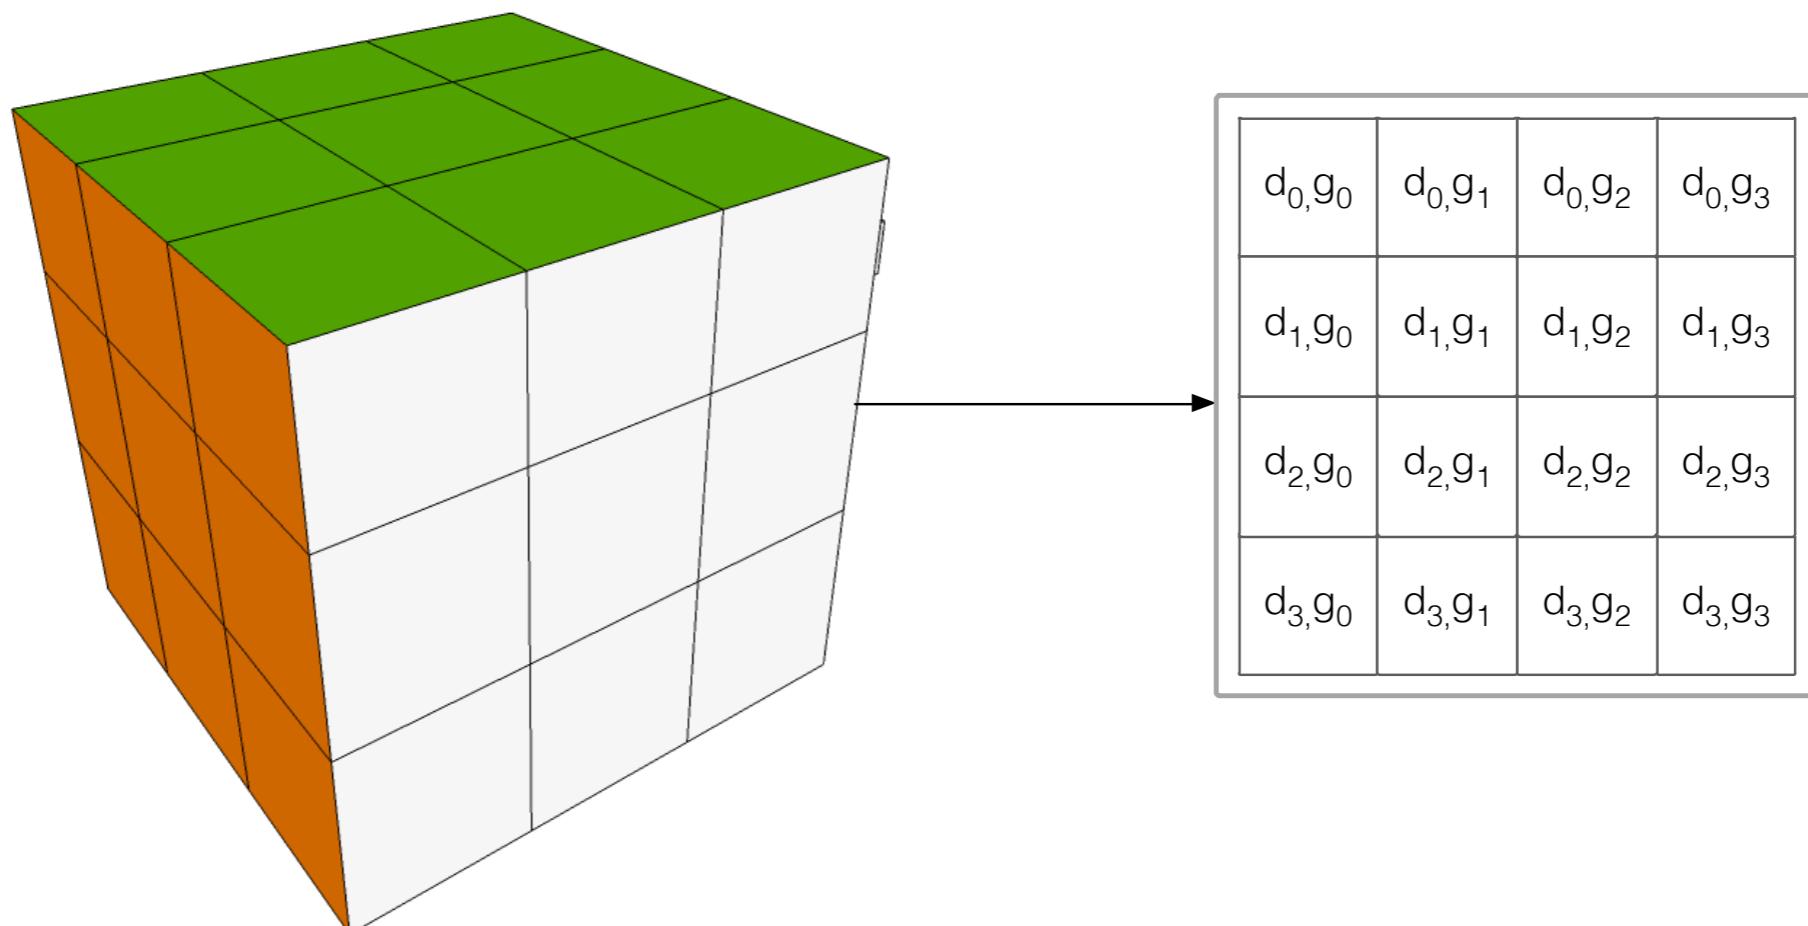


canonical order is i,j,k,d,g

```
make_slices_view<0,1,2>(  
    rhs  
)
```

(i, j, k)

(g,d)



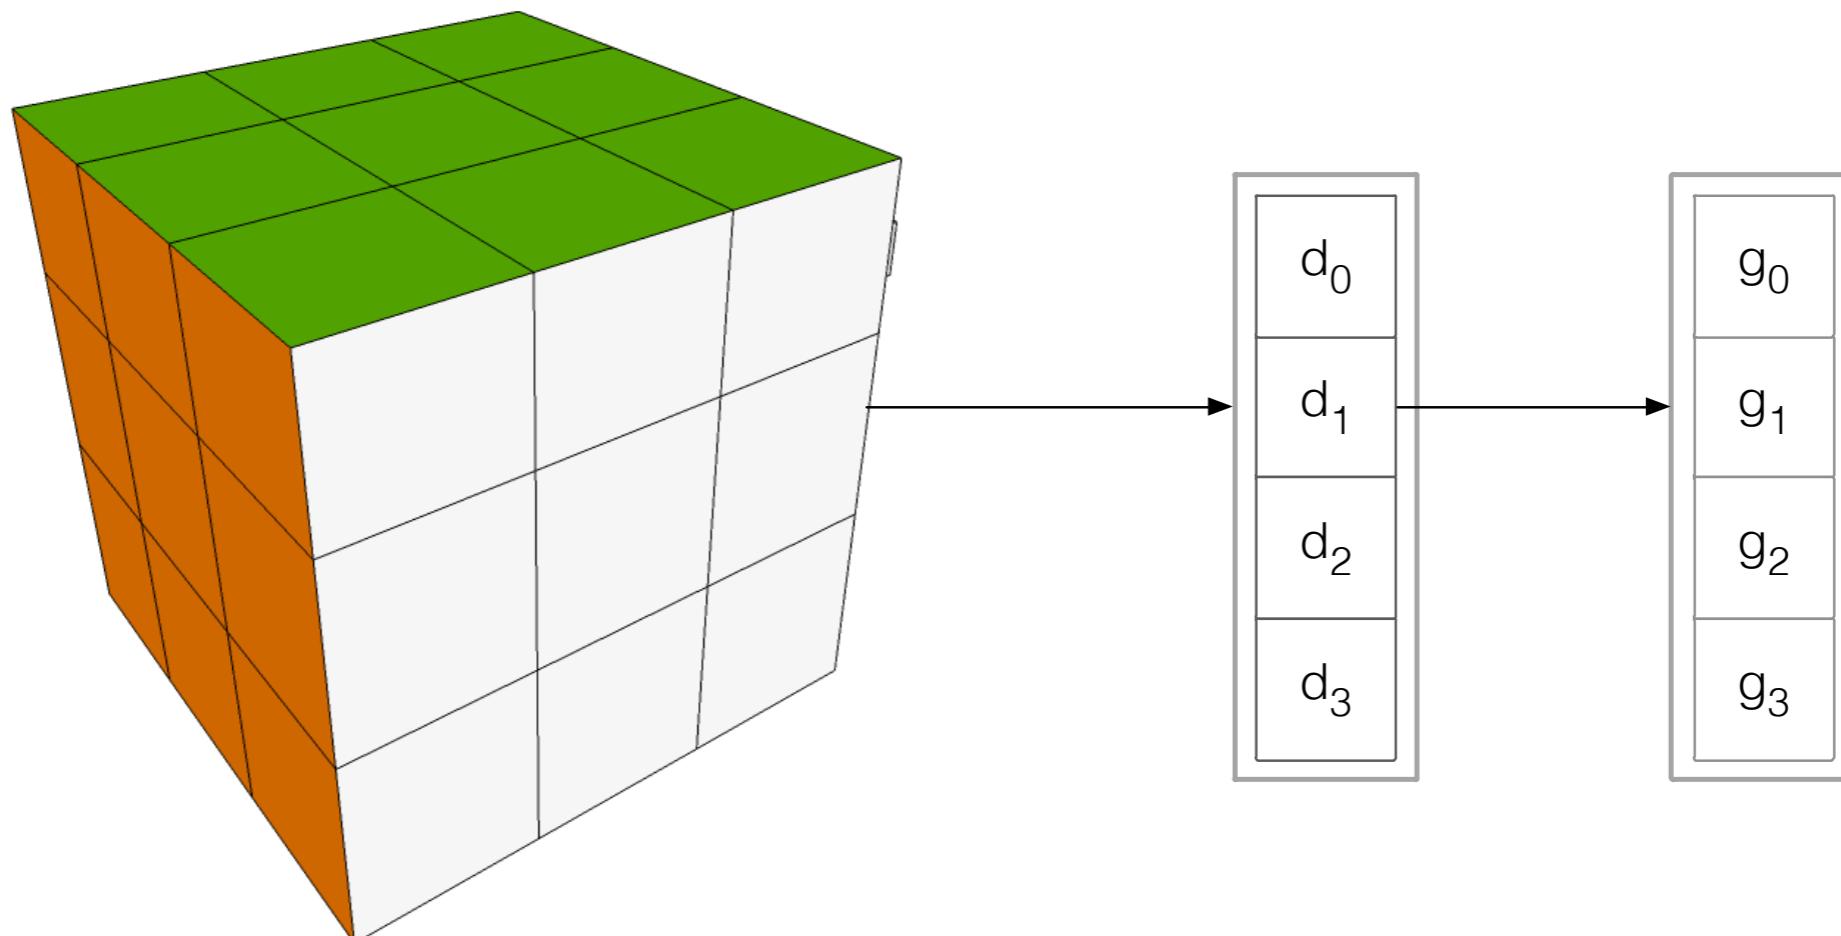
canonical order is i,j,k,d,g

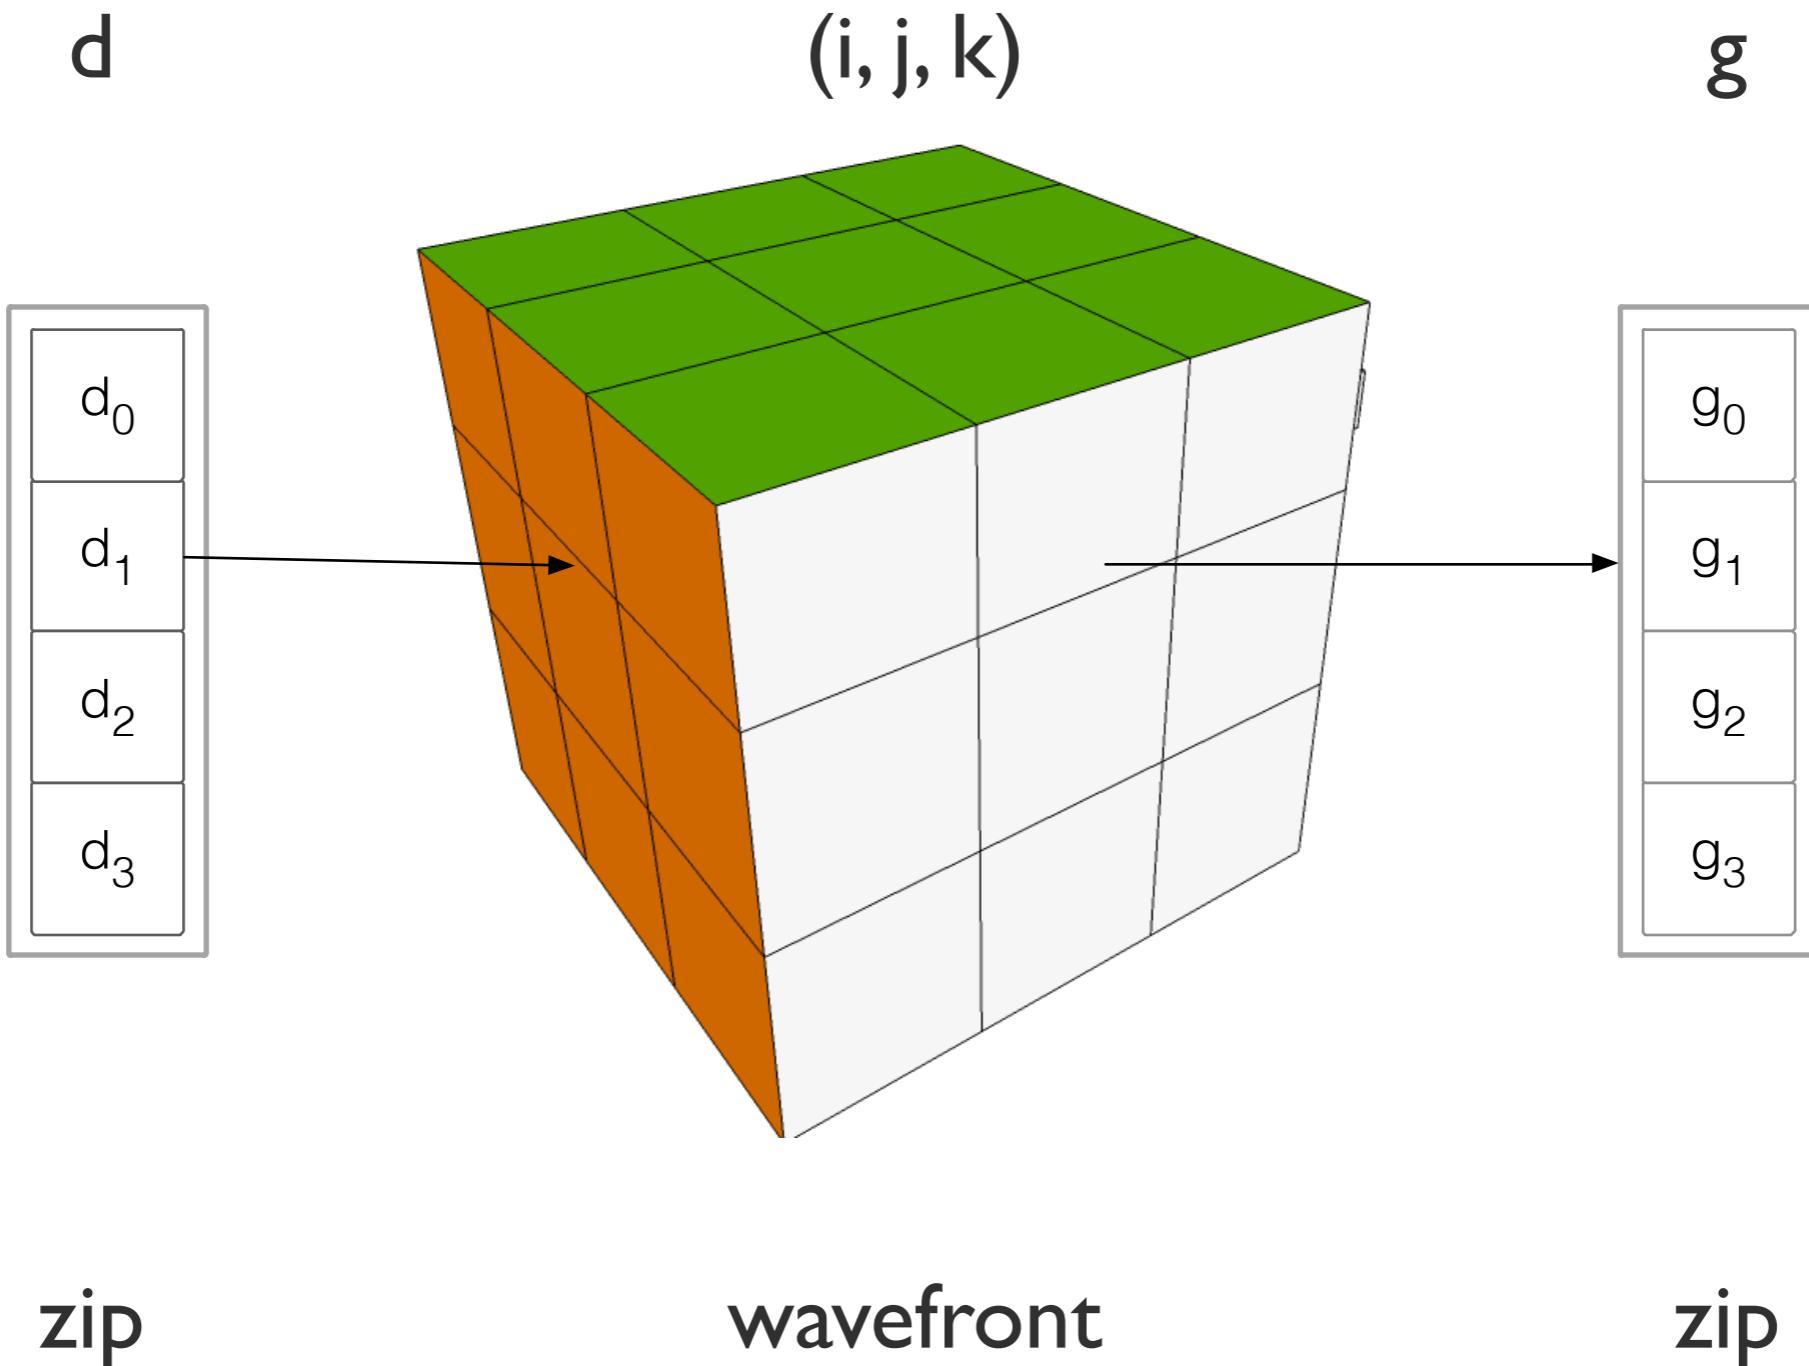
```
make_slices_view<0,1,2>(  
    make_slices_view<0,1,2,4>(rhs)  
)
```

(i, j, k)

g

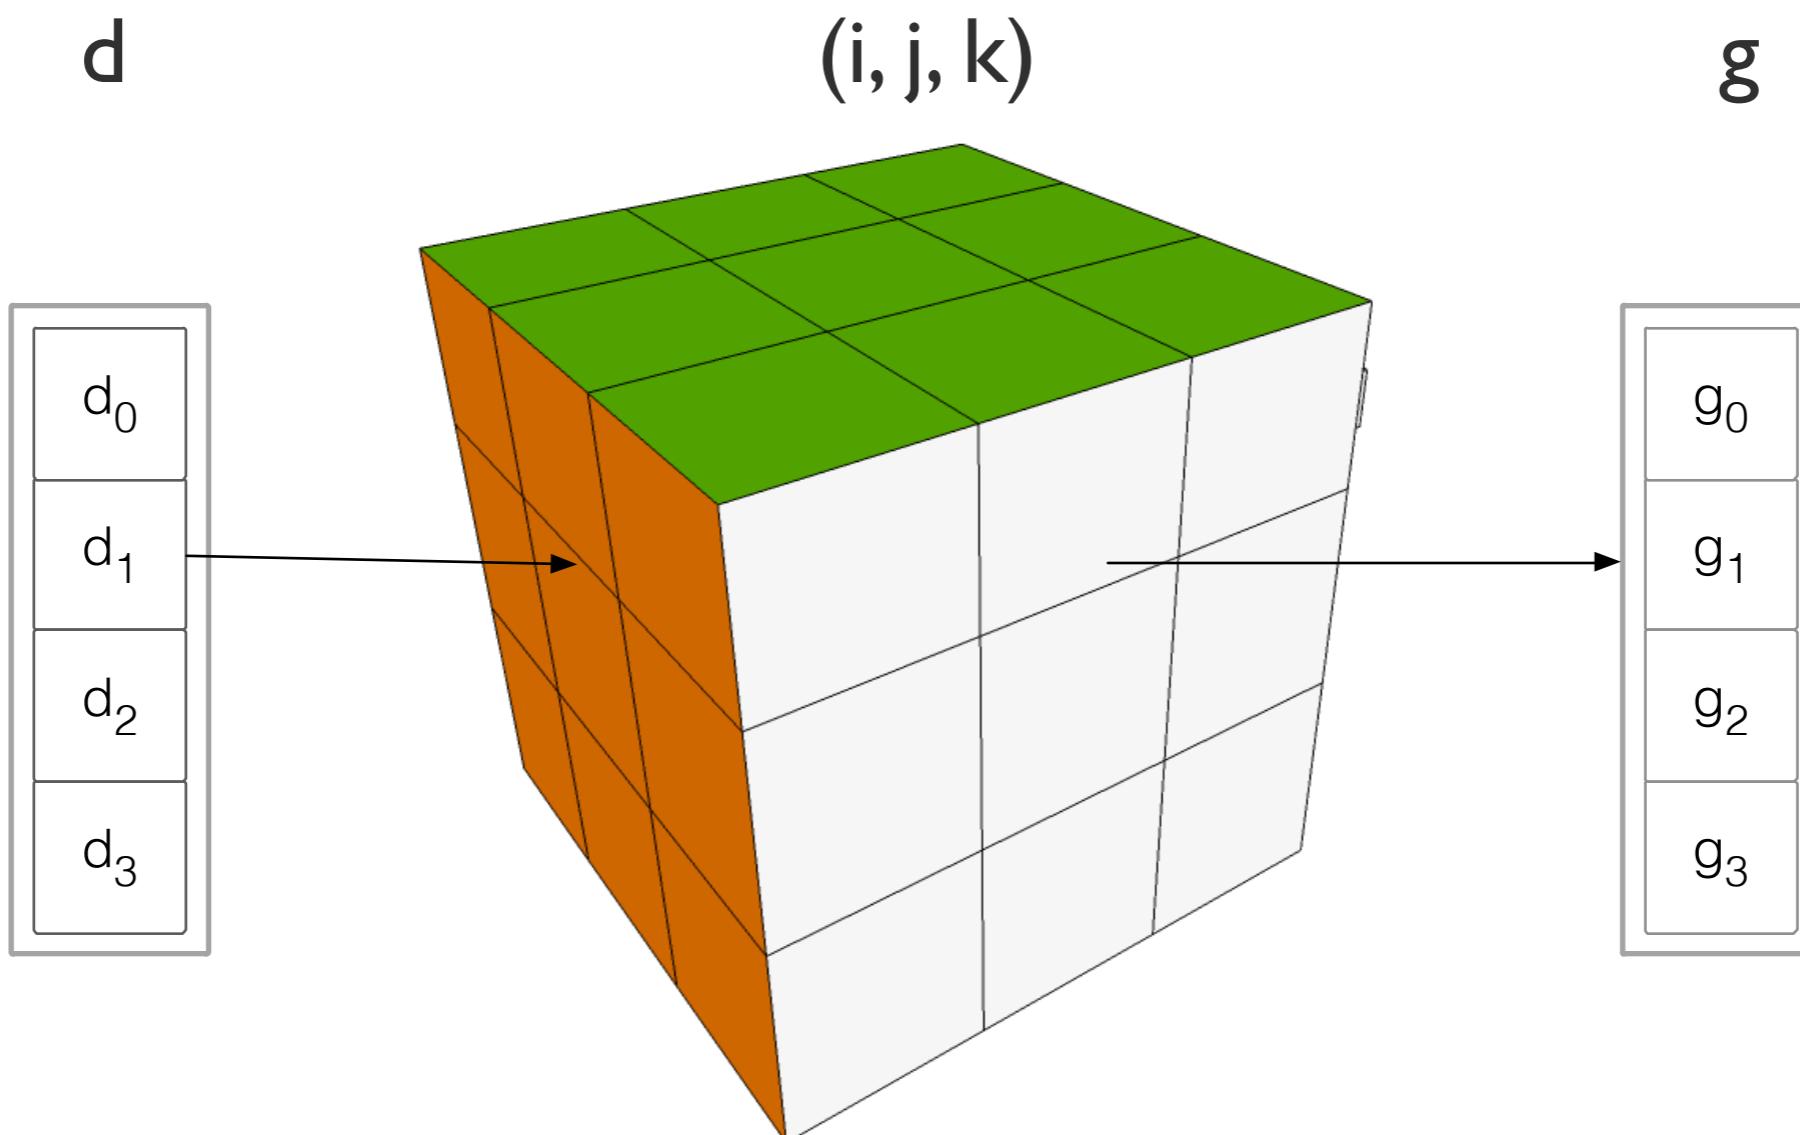
d





```
make_slices_view<3>(  
    make_slices_view<0,1,2,3>(rhs)  
)
```

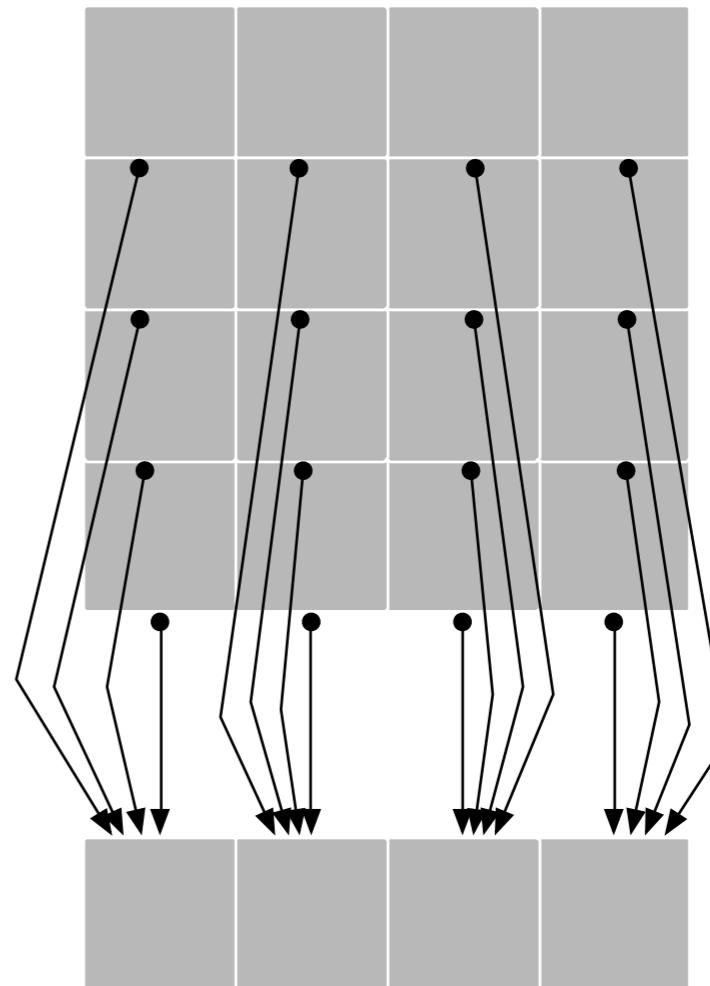
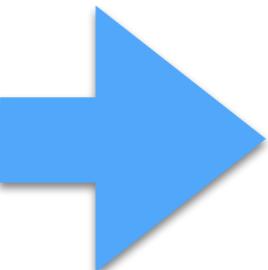
canonical order is i,j,k,d,g



- A view that extends a single dimension of a view to a given size
- Scalar expansion

```
make_extended_view<0>(vw, 4)
```

`multiarray<I>(4)`





Extended View in Kripke

- `sigt` is a container where the flux is aggregated for each direction
- The dimensions are (i,j,k,g) , but we want to extend the direction dimension d
 - Map 4D container to 5D view using expansion
- All views are then conformable to each other

```
ext_sigt = make_extended_view<3>(
    sigt, num_directions
);
skeleton(rhs, psi, ext_sigt)
```

- A view that evaluates a function when read
 - Collection of elements is closed-form solution

```
make_functor_view(tuple(4,4), f)
```

$f_{0,0}$	$f_{0,1}$	$f_{0,2}$	$f_{0,3}$
$f_{1,0}$	$f_{1,1}$	$f_{1,2}$	$f_{1,3}$
$f_{2,0}$	$f_{2,1}$	$f_{2,2}$	$f_{2,3}$
$f_{3,0}$	$f_{3,1}$	$f_{3,2}$	$f_{3,3}$



ZGD Kernel

```
skeleton = ...
containers = ...
```

```
execute(skeleton,
```

```
)
```



ZGD Kernel

```
skeleton = ...
```

```
containers = ...
```

```
ext_sigt = make_extended_view<3>(sigt, num_directions)
```

```
execute(skeleton,
```

```
)
```



ZGD Kernel

```
skeleton = ...
containers = ...

ext_sigt = make_extended_view<3>(sigt, num_directions)

psi_slices = make_slices_view<0,1,2>(
    make_slices_view<0,1,2,4>(psi)
)

rhs_slices = // same as above with rhs
ext_sigt_slices = // same as above with ext_sigt

execute(skeleton,
    psi_slices, rhs_slices, ext_sigt_slices, output
)
```



ZGD Kernel

```
skeleton = ...
containers = ...

ext_sigt = make_extended_view<3>(sigt, num_directions)

psi_slices = make_slices_view<0,1,2>(
    make_slices_view<0,1,2,4>(psi)
)

rhs_slices = // same as above with rhs
ext_sigt_slices = // same as above with ext_sigt

boundary_view0 = make_functor_view(dims, boundary_cond)

execute(skeleton,
    boundary_view0, boundary_view1, boundary_view2,
    psi_slices, rhs_slices, ext_sigt_slices, output
)
```



Implementation of views

- Create new views and modify components
- Views are defined as
 - C: partitioned collection of elements
 - D: domain of the view
 - F: function to map from the view to the container
 - O: operations



Implementation of views

- Slices view
 - Modify domain to include dimensions that are being sliced
 - Mapping function to project from n-d to m-d
- Extended view
 - Modify domain to include new dimension
 - Mapping function ignores indices in that dimension



- Kripke Overview
- Designing the STAPL implementation
 - Data structures
 - Skeleton based algorithms
 - Views
- Current work

- **Integration of multiarray<multiarray<>> for data structures.**
 - Uniform, spatial decomposition for outer container of zonesets.
 - Varying composition of subset of dimensions for inner container, based on specified kernel.
- **Integrated View composition for kernels**
 - slices_view and extended_view to conform data to kernel nesting.
- **Wavefront skeleton written and integrated both across zonesets and within a zoneset.**

- Skeleton Composition
- Performance Tuning
- Investigate further extensions, configurations.



Questions?